

# Decision Trees and ensemble methods

Pietro Gori

Associate Professor  
Equipe IMAGES - Télécom Paris  
[pietro.gori@telecom-paris.fr](mailto:pietro.gori@telecom-paris.fr)



# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Supervised Learning - Probabilistic framework

$X$ : input data  $x_i^j$ , random variable in  $\mathcal{X} = \mathbb{R}^p$  with  $i = 1, \dots, n$  and  $j = 1, \dots, p$  where  $n$  and  $p$  are the number of observations and variables respectively

$Y$ : response (to predict), random variable in  $\mathcal{Y} = \{C_1, \dots, C_K\}$  (classification with  $K$  classes) or  $\mathcal{Y} = \mathbb{R}$  (regression)

$P$ : joint probability distribution of  $(X, Y)$ , fixed but unknown

$\mathcal{D}_n = \{(\mathbf{x}_i, y_i) \in \mathcal{X} \times \mathcal{Y}\}$ : i.i.d. samples drawn from  $P$

$\mathcal{F}$ : collection of classifiers  $f \in \mathcal{F}$

$\mathcal{L}$ : loss function which measures the error of the classifier/model

- Examples (classification) :  $\mathcal{L}(\mathbf{x}, y, f(\mathbf{x})) = \begin{cases} 1, & \text{si } f(\mathbf{x}) \neq y, \\ 0, & \text{sinon.} \end{cases}$

- Example (regression):  $\mathcal{L}(\mathbf{x}, y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$

**Goal:** estimate from  $\mathcal{D}_n$  the function  $f \in \mathcal{F}$  which minimizes the risk (cost) function  $R(f) = \mathbb{E}_P[\mathcal{L}(X, Y, f(X))]$

# Estimate a classifier

We need to define:

- **input and output data space** ( $\mathcal{X} \mathcal{Y}$ )
- **type of classifier** ( $\mathcal{F}$ )
- **cost function** ( $\mathcal{L}$ ) to minimize for finding the best  $f$
- **minimization algorithm** for  $\mathcal{L}$
- **method for model selection** to estimate hyper-parameters
- **method to evaluate performance**

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees


- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Constant piecewise model/classifier

$\mathcal{F}$  belongs to the set of **constant piecewise functions**. We divide the input space  $\mathcal{X}$  in  $M$  disjoint partitions  $\mathcal{C}_m$ :  $\mathcal{X} = \mathcal{C}_1 \sqcup \dots \sqcup \mathcal{C}_M$ . To simplify things, we assume that the separation lines are parallel to the coordinate axes:



# Constant piecewise model/classifier

Let  $\mathbf{x} = (\mathbf{x}^1, \dots, \mathbf{x}^p)$ , we model the function  $f$  with a different constant value  $\alpha_m$  in every partition  $\mathcal{C}_m$ :  $f(\mathbf{x}) = \sum_{m=1}^M \alpha_m \mathbb{1}_{\mathbf{x} \in \mathcal{C}_m}$ , where  $\mathbb{1}_{\mathbf{x} \in \mathcal{C}_m}$  is equal to 1 if  $\mathbf{x} \in \mathcal{C}_m$  and 0 otherwise.

- **Regression:** If we use the L2-norm ( $\sum_i (y_i - f(\mathbf{x}_i))^2$ ), then the best  $\alpha_m^*$  is simply the average of the  $y_i$  within the region  $\mathcal{C}_m$ :  
$$\alpha_m^* = \frac{1}{|\mathcal{C}_m|} \sum_{\mathbf{x}_i \in \mathcal{C}_m} y_i, \text{ where } |\mathcal{C}_m| \text{ is the number of elements within } \mathcal{C}_m$$
- **Classification:** We define the proportion of observations belonging to class  $k$  within region  $\mathcal{C}_m$  as  $\rho_{mk} = \frac{1}{|\mathcal{C}_m|} \sum_{\mathbf{x}_i \in \mathcal{C}_m} \mathbb{1}_{y_i=k}$ . The best  $\alpha_m^*$  is then equal to the majority class:  $\alpha_m^* = \arg \max_k \rho_{mk}$



# Constant piecewise model/classifier

- Motivation: easy to interpret
- Limitations:
  - regions are difficult to describe
  - If the partition is fixed beforehand, many regions might end up being empty
- Possible solution: learn the partitions from the data ! How to avoid curse of dimensionality ?

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees



- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods



- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Decision trees

Presented by Breiman *et al.* in 1984 under the name of CART:  
Classification and Regression Trees.





# Decision trees



## First idea:

Use several hyperplanes (and not only one) to build non linear decision boundaries.



# Decision trees



## Secod idea:

Use separation lines orthogonal to the coordinate axis, *i.e.*, hyperplanes  $\{\mathbf{x} \in \mathcal{X} : \mathbf{x}^j = \tau\}$  to ease interpretation.

# Decision trees



## Third idea:

Use recursive binary decision trees: The full data-set sits at the top of the tree. Every node (junction) is associated to a separating hyperplane  $\{\mathbf{x} \in \mathcal{X} : \mathbf{x}^j = \tau\}$ . The terminal nodes correspond to the regions.

# Linear partition orthogonal to the axes

- We use the same model as before:  $f(\mathbf{x}) = \sum_{m=1}^M \alpha_m \mathbb{1}_{\mathbf{x} \in \mathcal{C}_m}$ , namely we assign a constant value at every region
- Let  $\mathbf{x} = (\mathbf{x}^1, \dots, \mathbf{x}^p)$  with  $p$  variables. We define the split  $t_{j,\tau}(\mathbf{x})$  along the direction  $\mathbf{x}^j$  with threshold  $\tau$  as:

$$t_{j,\tau}(\mathbf{x}) = \text{sign}(\mathbf{x}^j - \tau) = \begin{cases} +1, & \text{si } \mathbf{x}^j > \tau \\ -1, & \text{si } \mathbf{x}^j < \tau \end{cases} \quad (1)$$

- We assume for now that we have already defined a *loss function*  $\mathcal{L}$  and a *stopping criteria*

# Efficient recursive algorithm

For a binary tree:

- 1 Given  $\mathcal{D}_n$ , the entire data-set is the root node
- 2 Look for the best separator  $t_{j,\tau}$  on  $\mathcal{D}_n$  such that the local cost function  $\mathcal{L}(t_{j,\tau}, \mathcal{D}_n)$  is minimal. This means looking for the “best” direction  $j$  and threshold  $\tau$  that splits  $\mathcal{D}_n$  into  $\mathcal{D}_n^d$  and  $\mathcal{D}_n^g$ .
- 3 Be careful: the splitting values  $\tau$  are not infinite ! They depend on the data  $\mathcal{D}_n$ . Hence, we can scan all inputs  $\mathbf{x}$  and quickly find the best  $j$  and  $\tau$
- 4 Split  $\mathcal{D}_n$  to  $\mathcal{D}_n^d$  and  $\mathcal{D}_n^g$  using the estimated separator hyperplane.
- 5 It results two nodes, a left ( $\mathcal{D}_n^g$ ) and a right ( $\mathcal{D}_n^d$ ) one
- 6 Evaluate the stopping criteria for the right node, if it is verified, the nodes becomes a terminal node, otherwise go to 3 using  $\mathcal{D}_n^d$  as input space
- 7 Evaluate the stopping criteria for the left node, if it is verified, the nodes becomes a terminal node, otherwise go to 3 using  $\mathcal{D}_n^g$  as input space



# Examples

Given the input data-set  $\mathcal{D}_n$  and a binary separator  $t_{j,\tau}$ , we have

$$\mathcal{D}_n^d(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) > 0\}$$


$$\mathcal{D}_n^g(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) \leq 0\}$$

# Examples

Given the input data-set  $\mathcal{D}_n$  and a binary separator  $t_{j,\tau}$ , we have

$$\mathcal{D}_n^d(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) > 0\}$$

$$\mathcal{D}_n^g(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) \leq 0\}$$





# Examples


Given the input data-set  $\mathcal{D}_n$  and a binary separator  $t_{j,\tau}$ , we have

$$\mathcal{D}_n^d(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) > 0\}$$


$$\mathcal{D}_n^g(j, \tau) = \{(\mathbf{x}, y) \in \mathcal{D}_n, t_{j,\tau}(\mathbf{x}) \leq 0\}$$




# Example



# Example



# Example



- The presented algorithm (called CART) is greedy
- We do not optimize a global criterion. Instead, we locally look for an optimal separator (with respect to  $L$ ), which means at every direction  $j$  independently.
- Why not using a global criterion ? Why not looking for the best sequence of separators ?

- The presented algorithm (called CART) is greedy
- We do not optimize a global criterion. Instead, we locally look for an optimal separator (with respect to  $L$ ), which means at every direction  $j$  independently.
- Why not using a global criterion ? Why not looking for the best sequence of separators ? Because it is a NP-hard problem !



# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

- We recall that for a classification tree, given the input data  $\mathcal{D}_n$  divided in  $K$  classes, we define the proportion of observations belonging to class  $k$  as:  $\rho_k(\mathcal{D}_n) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}(y_i = k)$
- Note that  $\rho(\mathcal{D}_n) = (\rho_1(\mathcal{D}_n), \dots, \rho_K(\mathcal{D}_n))^T \in \Delta_{K-1}$  where 
$$\Delta_{K-1} := \left\{ \rho_k \in \mathbb{R}^K : \sum_{k=1}^K \rho_k = 1 \text{ and } \forall k \in \llbracket 1, K \rrbracket, \rho_k \geq 0 \right\}$$
 is the  $(K-1)$ -simplex. A simplex can be seen as the smallest convex set containing the given vertices, or the convex hull of the  $K$  vertices.  
*Example:* a 2-simplex is a triangle.

# Cost function - Classification

Among all parameters  $(j, \tau) \in \{1, \dots, p\} \times \{\tau_1, \dots, \tau_m\}$ , we look for  $j^*$  and  $\tau^*$  which minimizes the following cost function:

$$L(t_{j,\tau}, \mathcal{D}_n) = \frac{n_g}{n} H(\rho(\mathcal{D}_n^g(j, \tau))) + \frac{n_d}{n} H(\rho(\mathcal{D}_n^d(j, \tau)))$$

avec  $n_g = |\mathcal{D}_n^g(j, \tau)|$  et  $n_d = |\mathcal{D}_n^d(j, \tau)|$

- $H$  is an “impurity” function that evaluate the splitting. Pure means a node with observations from the same class. We want to **minimize**  $H$  in order to have pure nodes.
- The total cost is the sum of the impurity of each child node ( $\mathcal{D}_n^g$  and  $\mathcal{D}_n^d$ ) weighted by the proportion of its observations ( $n_g$  and  $n_d$ )
- We evaluate a finite number of thresholds (max  $n$ )

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- **Impurity function**
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

## Definition: impurity function

An impurity function  $H : [0 : 1]^K \rightarrow \mathbb{R}$  is a function defined on  $\Delta_{K-1}$  for which the following properties hold:

- ①  $H$  becomes maximum at points  $(\frac{1}{K}, \dots, \frac{1}{K})^\top$ , i.e., all  $\rho_k$  are equal
- ②  $H$  becomes minimum at points  $(1, 0, \dots, 0)^\top, (0, 1, 0, \dots, 0)^\top, \dots, (0, \dots, 0, 1)^\top$ , i.e., the probability of being in a certain class is 1 and 0 for all other classes. These are the vertices of  $\Delta_{K-1}$ .
- ③  $H$  is symmetric with respect to its arguments  $\rho_1, \dots, \rho_K$ , i.e., even if we permute  $\rho_j$ ,  $H$  does not change

# Impurity function: binary case ( $K = 2$ )

When we have only two classes ( $K = 2$ ):

- $\Delta_{K-1}$  is the line segment joining  $(1, 0)$  and  $(0, 1)$  in  $\mathbb{R}^2$
- $H$  becomes maximum at  $(\frac{1}{2}, \frac{1}{2})$
- $H$  becomes minimum at  $(0, 1)$  or  $(1, 0)$ , which means when all observations in one region belong to the same class.

# Misclassification error

Given the data of a node  $\mathcal{D}_n$  (it might be the root node or a child node), we assign the observations in  $\mathcal{D}_n$  to the majority class  $k^*$ :

$$k^* = \arg \max_{k=1,\dots,K} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_n} \mathbb{1}(y_i = k)$$

Then we define:


$$\text{Misclassification error: } H_{\text{mis}}(\mathcal{D}_n) = 1 - \rho_{k^*}(\mathcal{D}_n),$$

This is the error that we commit by assigning the observations to the class  $k^*$ . Remember that  $\sum_{k=1}^K \rho_k = 1$ .

# Misclassification error

When the number of classes  $K$  is 2

$$H_{\text{mis}}(\mathcal{D}_n) = 1 - \max_{k=1,2} \rho_k(\mathcal{D}_n) = \min(\rho_1(\mathcal{D}_n), 1 - \rho_1(\mathcal{D}_n))$$






# Limitations of the misclassification error

- Remember that the cost of a split is:


$$L(t_{j,\tau}, \mathcal{D}_n) = \frac{n_g}{n} H(\rho(\mathcal{D}_n^g(j, \tau))) + \frac{n_d}{n} H(\rho(\mathcal{D}_n^d(j, \tau)))$$

- For both splits we have  $L_{\text{mis}} = \frac{1}{2} \cdot \frac{1}{4} + \frac{1}{2} \cdot \frac{1}{4} = \frac{1}{4} = \frac{3}{4} \cdot \frac{1}{3} + \frac{1}{4} \cdot 0 = \frac{1}{4}$ .  
However it seems that the second one is definitely better !



# Limitations of the misclassification error

- For a partition where a class has a clear majority, we might not find a split which reduces  $L$
- The function is not differentiable (optimization is harder)
- It might underestimate pure nodes:



# Strict impurity function

## Definition: strict impurity function

Let  $H : [0 : 1]^K \rightarrow \mathbb{R}$  be an impurity function,  $\rho, \rho'$  two distributions in  $\Delta_{K-1}$  with  $\rho \neq \rho'$  and  $\alpha \in ]0, 1[$ . Then  $H$  is called strict, if it is strictly concave:

$$H(\alpha\rho + (1 - \alpha)\rho') > \alpha H(\rho) + (1 - \alpha)H(\rho')$$

If  $H$  is strict then it follows that

$$L(t_{j,\tau}, \mathcal{D}_n) = \frac{n_g}{n} H(\rho(\mathcal{D}_n^g(j, \tau))) + \frac{n_d}{n} H(\rho(\mathcal{D}_n^d(j, \tau))) \leq H(\rho(\mathcal{D}_n))$$
$$n_g = |\mathcal{D}_n^g(j, \tau)| \quad \text{et} \quad n_d = |\mathcal{D}_n^d(j, \tau)|$$

the equality is given iff  $\rho_k(\mathcal{D}_n) = \rho_k(\mathcal{D}_n^g) = \rho_k(\mathcal{D}_n^d)$  for all  $k$

Remark: The impurity function of the misclassification error is concave, but it is not strictly concave.  $L$  might be equal for all possible splittings.


$$\text{Entropy: } H_{\text{ent}}(\mathcal{D}_n) = - \sum_{k=1}^K \rho_k(\mathcal{D}_n) \log \rho_k(\mathcal{D}_n)$$

- if we use  $\log_2$ , it is called Shannon entropy
- $-\log \rho(\mathcal{D}_n)$  is the information content of  $\mathcal{D}_n$
- Entropy is defined as the expected value of the information content (average amount of information). It measures the randomness
- when an event is certain, entropy is 0
- information gain is defined as reduction in entropy
- it is differentiable


# Entropy

When the number of classes  $K$  is 2

$$H_{\text{ent}}(\mathcal{D}_n) = -\rho_1(\mathcal{D}_n) \log(\rho_1(\mathcal{D}_n)) - (1 - \rho_1(\mathcal{D}_n)) \log(1 - \rho_1(\mathcal{D}_n))$$



# Example



Question: Compute  $L_{\text{ent}}$  associated to  $H_{\text{ent}}$ . Which split is better ?

$$H_{\text{Gini}}(\mathcal{D}_n) = \sum_{k=1}^K \rho_k(\mathcal{D}_n)(1 - \rho_k(\mathcal{D}_n)) = \sum_{k=1}^K \sum_{\substack{k'=1 \\ k' \neq k}}^K \rho_k(\mathcal{D}_n)\rho_{k'}(\mathcal{D}_n)$$


Two different interpretations:

- If we code each observation as 1 for class  $k$  and zero otherwise (bernoulli variable) the variance is  $\rho_k(\mathcal{D}_n)(1 - \rho_k(\mathcal{D}_n))$ . The Gini index is the sum of the variances of the “binarized” classes
- We do not assign observations to the majority class (as for  $H_{\text{mis}}$ ) but we classify them to class  $k$  with probability  $\rho_k(\mathcal{D}_n)$ . The training error rate of this rule is  $\sum_{\substack{k'=1 \\ k' \neq k}}^K \rho_k(\mathcal{D}_n)\rho_{k'}(\mathcal{D}_n)$ . The Gini index is the sum over all classes.

# Gini index


When the number of classes  $K$  is 2

$$H_{\text{Gini}}(\mathcal{D}_n) = 2\rho_1(\mathcal{D}_n) (1 - \rho_1(\mathcal{D}_n))$$





# Example



Question: Compute  $L_{\text{Gini}}$  associated to  $H_{\text{Gini}}$ . Which split is better ?

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Stopping criteria

Without stopping criteria, we could grow a tree until a situation where each observation represents a terminal node. This would be computationally expensive, difficult to interpret and prone to over-fitting. Instead, we could use one or more of the following stopping criteria:

- maximal depth
- maximum number of terminal nodes
- a node becomes a terminal node when it reaches a maximum number of observations
- degree of purity of a node (*i.e.*, threshold on  $\rho_k(\mathcal{D}_n)$ )

- For a binary tree : if we have a categorical variable  $x$  which can take up to  $M$  values, we transform it into  $M$  binary variables
- Warning: The partitioning algorithm tends to favor categorical variables with many values since the number of possible partitions grows exponentially with  $M$ . This means that we have more choices to find a good partition. This can lead to over-fitting ! Try to avoid such variables.

# Loss matrix

In some cases, the consequences of misclassifying observations can be very serious (*i.e.*, medicine). To account for that, we introduce a loss matrix  $C \in \mathbb{R}^{K \times K}$ , with  $C_{k,k'}$  being the loss incurred for classifying observations of class  $k$  as belonging to class  $k'$

$$C_{k,k'} = 0 \text{ si } k = k' \quad C_{k,k'} \geq 0 \text{ si } k \neq k'$$

We can then modify the Gini index as follows:

$$\text{Gini index: } \sum_{k=1}^K \sum_{\substack{k'=1 \\ k' \neq k}}^K C_{k,k'} \rho_k(\mathcal{D}_n) \rho_{k'}(\mathcal{D}_n)$$

Note: This works for  $K > 2$  but it has no effect in the binary case (Why?).

# Regression trees

For regression the process is almost identical, we only change the impurity function. We use the squared error (or variance):

$$H(\mathcal{D}_n) = \frac{1}{|\mathcal{D}_n|} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_n} (y_i - \bar{y}_n)^2$$

where

$$\bar{y}_n = \frac{1}{|\mathcal{D}_n|} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_n} y_i$$

and we minimize as before

$$L(t_{j,\tau}, \mathcal{D}_n) = \frac{n_g}{n} H(\mathcal{D}_n^g(j, \tau)) + \frac{n_d}{n} H(\mathcal{D}_n^d(j, \tau))$$

Note: as before we want to maximize the homogeneity (purity) of the terminal nodes

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- **Model selection**

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Model selection (1)

We can compute one (or more) of the following hyper-parameters instead than fixing them as stopping criteria:

- maximal depth of the tree
- maximum number of terminal nodes
- maximum number of observations in a node to become a terminal node

→ we could use cross validation



## Pruning (2)

What's the optimal size of a tree ? A large tree might overfit the data, while a small tree might not capture important structures.

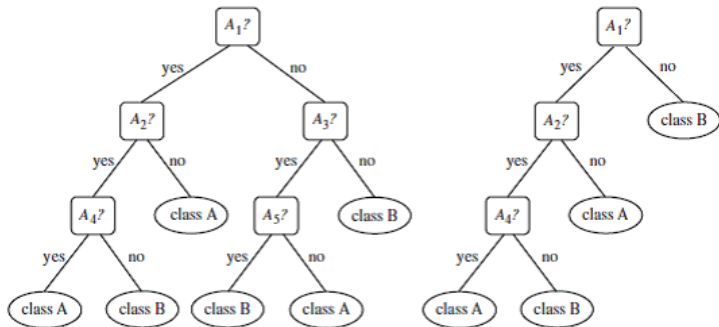
A possible solution is to grow a large tree in a training set stopping the splitting only when the terminal nodes have reached a minimum number of observations or a certain purity. Then, we produce several new trees by pruning the original tree at different nodes. When we want to prune with respect to node  $t$ , we delete all successor nodes of  $t$  in the original tree.

The new trees are then tested in a validation set. We select the tree that gives the best performance.

Trying all possible trees might be computationally unfeasible. Several greedy techniques exist. See Hastie *et al.* (2009) for more details.

Note: Only *Minimal Cost-Complexity Pruning* is currently supported in `sklearn`

# Example of pruning



## Advantages

- Build a non-linear and interpretable decision function
- Invariant under scaling and other linear transformations of the input data  $X$
- Robust to the inclusion of (few) irrelevant features  $x$
- It works for multi-class
- Computationally efficient:  $O(\log F)$ , where  $F$  is the number of terminal nodes
- It works for continuous and categorical variables (even if not optimal)

## Drawbacks

- Low bias but very high variance. A small change in (all) input data can bring to a completely different tree (noisy) ! This instability is due to the hierarchical nature of the process. → averaging trees reduces the variance (bagging, random forests)
- No global optimization (NP-complete) we use a greedy algorithm where locally optimal decisions are made
- Separation hyperplanes are aligned with the feature axes → it might entail a sub-optimal solution
- Splits are hard. This creates piecewise-constant predictions with discontinuities at the split boundaries → prediction function is not smooth !

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods


- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Trees limitations

- We know that Trees suffer from high variance
- We also know that averaging reduce variance  $\rightarrow$  given  $M$  iid observations, each with variance  $\sigma^2$ , the variance of the average is  $\frac{\sigma^2}{M}$
- For any statistical model (specially trees) we could take many different training sets, build a separate model in each training set and average the resulting predictions for regression or using the majority vote for classification
- Problem: we usually have only one training set !!  $\rightarrow$  **Ensemble methods**

# Qui veut gagner des millions ?

Which answer is the best choice ?



# Qui veut gagner des millions ?

- Assume that you first use the option *Remove two incorrect answers* removing answer A) and D) → binary classification problem !
- Assume that:
  - ① The number of people in the audience  $N$  is **odd** (to have only one winner...)
  - ② The probability of choosing the correct answer is  $p$  for every person
  - ③ The votes of the audience are **independent**
- The probability of getting exactly  $k$  correct answers among the  $N$  voters is :  $\binom{N}{k} p^k (1-p)^{N-k}$  (binomial distribution)
- The probability of ensemble success, namely when the majority  $(N/2) + 1$  makes the correct decision, is :
$$\sum_{k=(N/2)+1}^N \binom{N}{k} p^k (1-p)^{N-k}$$
- This sum approaches 1 as  $N \rightarrow \infty$  only when  $p > 0.5$ , otherwise it tends towards 0



# Qui veut gagner des millions ?

## Condorcet's jury Theorem

Given a group of  $N$  independent voters who choose the correct answer with probability  $p > 0.5$  then, as  $N \rightarrow \infty$ , the probability of choosing the correct answer by majority vote tends to 1

- Assumptions about independence and uniform probability  $p$  not always realistic
- It can be extended beyond binary classification problem [1]
- It shows that **combining several weak classifiers may be better than using a single-expert classifier**

# Ensemble learning

## Ensemble learning - Definition

Combine several simple and moderately inaccurate classifiers to create a very accurate one

## Rationale

It may be easier and more accurate to train and combine several simple classifiers than to learn a single complex classifier

## Three famous ensemble methods

- 1 **Bagging**, 1994, Breiman
- 2 **Random Forests**, 1995, Tin Kam Ho
- 3 **Boosting**, 1990-1997, Schapire and Freund (AdaBoost won the Godel Prize in 2003)

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Bagging (Bootstrap aggregating)

- Aggregate the prediction of several weak classifiers fitted on randomly generated and (almost) independent training sets (*model averaging*) in order to reduce variance and avoid overfitting
- Given the training set  $\mathcal{D}_n$ , we generate  $M$  new training sets  $\{\mathcal{D}_k\}_1^M$ , each of size  $k$ , by randomly sampling with replacement  $k$  samples from  $\mathcal{D}_n$ . This is called **bootstrap** and it means that an observation  $\{x_i, y_i\}$  may appear multiple times in the same bootstrap training set  $\mathcal{D}_k$ .
- **Hypotheses**
  - 1 *Representativity*: if the size of  $\mathcal{D}_n$  is large enough to capture the underlying actual distribution of the data, then a bootstrap set  $\mathcal{D}_k$  should be a good approximation
  - 2 *Independence*: if  $k$  is sufficiently smaller than  $n$ , then the bootstrap sets  $\mathcal{D}_k$  should not be too much correlated among them

# Bagging (Bootstrap aggregating)

- Just like averaging i.i.d. random variables preserves the expected value and reduce variance, aggregating the prediction of several weak learners trained on (almost) independent data-sets does not change the expected outcome but reduces its variance
- Once fitted  $M$  weak models  $f_i$ , we can aggregate their predictions using:
  - A simple average for regression problems  $\frac{1}{M} \sum_{i=1}^M f_i$
  - A majority vote for classification problems
- **Possible limitations:**
  - If  $n$  is not big enough and  $k < n$ , the bootstrap sets will not be a good approximation of the actual underlying distribution of the data
  - if  $k \sim n$  the bootstrap sets might not be independent (correlation between predictions  $\rightarrow$  Random Forests)

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking

# Random Forests

- When  $k \sim n$ , bagged trees might be very similar to each other if there is one or more strong predictors (i.e. features)  $\rightarrow$  all trees will use these features in the top split
- If trees are similar their prediction might be highly correlated !  $\rightarrow$  Average of  $M$  identically distributed variables, each with variance  $\sigma^2$  and positive pairwise correlation  $\rho$ , has variance:

$$\rho\sigma^2 + \frac{1-\rho}{M}\sigma^2 \quad (2)$$

- We should create **many uncorrelated** bagged trees in order to reduce both terms  $\rightarrow$  **Random Forests** !
- The idea of Random Forests is to consider only a subsample of the  $p$  features at each split  $\rightarrow$  in many splits strong predictors will not be considered and thus bagged trees should be less correlated

# Random Forests - Algorithm

- For each one of the  $M$  tree  $T_m$ :
  - ① Draw a bootstrap sample from  $\mathcal{D}_n$
  - ② Until every terminal node has less than  $n_{min}$  observations (typical values are 1 for classification and 5 for regression), repeat the following steps for each terminal node:
    - ① Select at random  $q$  variables from the  $p$  of  $\mathbf{x}$  ( $q \leq p$ ). Typical values are  $\sqrt{p}$  for classification and  $p/3$  for regression
    - ② Estimate the best split for every selected variable
    - ③ Pick the best (variable,split) among the  $q$
    - ④ Split the node into left and right daughter nodes
- To predict the output of a new test observation  $\mathbf{x}$ :
  - *Regression*:  $f_{RF}^* = \frac{1}{M} \sum_{m=1}^M T_m(\mathbf{x})$
  - *Classification*: We use the majority vote among the classed predicted by each tree  $T_m(\mathbf{x})$
- The hyperparameters here are  $M$ ,  $q$  and  $n_{min} \rightarrow$  Cross-validation
- Intuitively, reducing  $q$  will also decrease  $\rho$  between pair of trees  $\rightarrow$  true only if the fraction of relevant variables is not small !



# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier


## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- **Boosting**
- Stacking

- Boosting is another general approach whose goal is to boost a weak learning algorithm, that works just slightly better than random guess, into a strong one
- There are many boosting algorithms and among them two of the most famous ones are:
  - 1 Adaptive boosting (AdaBoost)
  - 2 Gradient boosting (three of the most famous tree-based implementations are XGBoost, CatBoost and LightGBM)
- The first one (AdaBoost) was proposed by Schapire in 1990 and it consists of applying the same weak algorithm (e.g. Trees) several times, each time to a *modified version of the data*, producing a sequence of weak classifiers which are *finally combined*



**Figure 1:** Schematic illustration of the AdaBoost algorithm. Each weak classifier  $f_m$  is trained on a weighted variation  $w_n^{(m)}$  of the data. Classifiers are then combined into  $F_M$  via a weighted sum ( $\alpha_m$ ). Image modified from [2].

**Input:**  $\mathcal{D}_N = \{(x_i, y_i)\}$ , type of classifier  $f(x)$ , number of classifiers  $M$

---

- Initialize the weight of the data for the first classifier as  $1/N$ , namely all data points  $\{x_i, y_i\}$  have equal weight  $w_i^1 = 1/N$
  - For  $m = 1, \dots, M$ 
    - ① Fit classifier  $f_m$  to the training data  $\mathcal{D}_N$  weighted with  $\{w_i^m\}$ , namely minimize  $\sum_{i=1}^N w_i^m I(y_i \neq f_m(x_i))$
    - ② Define  $\epsilon_m = \frac{\sum_{i=1}^N w_i^m I(y_i \neq f_m(x_i))}{\sum_{i=1}^N w_i^m}$
    - ③ Use  $\epsilon_m$  to evaluate the weight  $\alpha_m$  of the classifier  $f_m \rightarrow$  The more accurate  $f_m$ , the greater its weight  $\alpha_m$
    - ④ Use  $\epsilon_m$  to update the weights of the data  $w_i^{m+1} \rightarrow$  higher weights are given to the data which were incorrectly predicted
- 

**Output:**  $F_M(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m f_m(x) \right)$

How are defined the weights from a mathematical point of view ?

- The weights of the classifiers:  $\alpha_m = \ln \left( \frac{1-\epsilon_m}{\epsilon_m} \right)$ . Thus  $\alpha_m$  takes value between 0 (classifier is equal to random guess) and  $+\infty$  (classifier is correct  $\forall i$ )
- For the weights of the data we have  $w_i^{m+1} = w_i^m \exp(\alpha_m)$  if  $y_i \neq f_m(x_i)$  and  $w_i^{m+1} = w_i^m$  if  $y_i = f_m(x_i)$ . Thus, since  $\alpha_m \in \mathcal{R}_+$ ,  $w_i^{m+1}$  will be equal to  $w_i^m$  if the classifier is correct and **bigger** if the classifier is wrong
- Please note that we could also choose to reduce the weights  $w_i^{m+1}$  for the observations correctly classified by defining  $w_i^{m+1} = w_i^m \exp(-\alpha_m)$  if  $y_i = f_m(x_i)$

# AdaBoost - binary classification

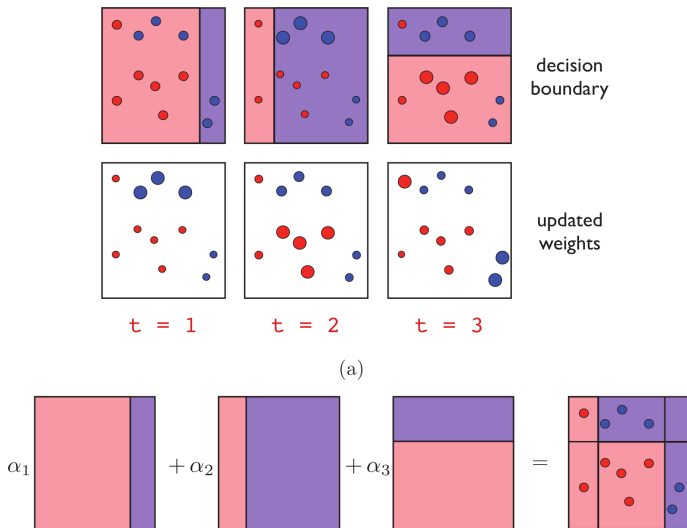
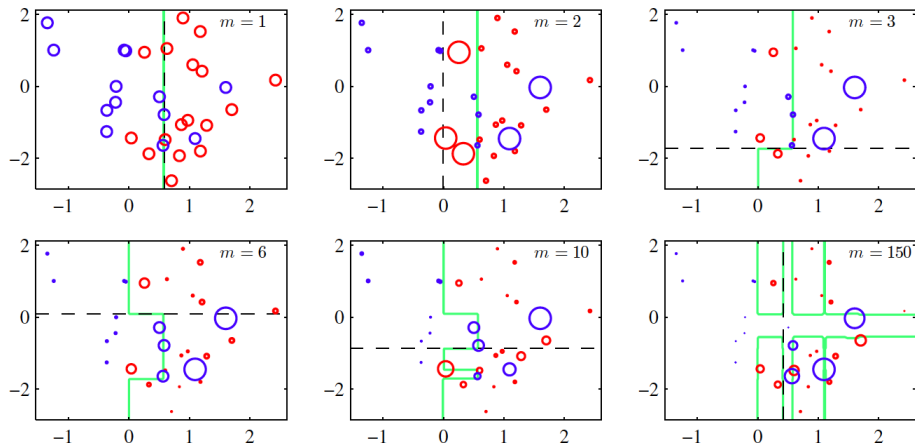


Figure 2: Axis-aligned hyperplanes as base learners. Image taken from [7].

# AdaBoost - binary classification



**Figure 3:** Black line: base learner at current iteration. Green line: ensemble decision boundary. Image taken from [2].

# Gradient Boosting

- To introduce gradient boosting it is useful to rewrite the AdaBoost algorithm in terms of *Forward Stagewise Additive Modeling*
- We can notice that the idea of Boosting is to add at each iteration a new weak classifier (or model)  $f(x; \beta_m)$  without adjusting the parameters  $\beta$  and coefficients  $\alpha$  of the previous models  $(1, .., m - 1)$   
→ greedy solution

$$h(x) = \frac{1}{2} \sum_{m=1}^M \alpha_m f(x; \beta_m) \quad (3)$$

- Please note that in this example  $f$  is a tree model,  $\beta_m$  the parameters of the  $m$  model (split variables, split points, predictions of terminal nodes) and  $\alpha_m$  are the coefficients of the  $m$  model



# Gradient Boosting

- At each iteration  $m$  one looks for the optimal model  $f_m = f(x; \beta_m)$  and corresponding coefficient  $\alpha_m$  to add to the current expansion  $h_{m-1}(x)$
- Given  $\mathcal{D}_N = \{(x_i, y_i)\}$ , number of models  $M$  and a loss function  $\mathcal{L}$
- The algorithm for the Forward Stagewise Additive Modeling is:

- 
- Initialize  $h_0(x) = 0$
  - For  $m = 1, \dots, M$ 
    - ① Compute

$$(\alpha_m, \beta_m) = \arg \min_{\alpha, \beta} \sum_{i=1}^N \mathcal{L}(y_i, h_{m-1}(x_i) + \frac{1}{2} \alpha f(x_i; \beta)) \quad (4)$$

- ② Set  $h_m(x) = h_{m-1} + \frac{1}{2} \alpha_m f(x; \beta_m)$

- This can be seen as a generic formulation for Boosting  $\rightarrow$  change the loss  $\mathcal{L}$  and you have different algorithms !
- Exponential loss ( $\mathcal{L}(y, f(x)) = \exp(-yf(x))$ )  $\rightarrow$  AdaBoost
- L2-norm, L1-norm, M-estimators  $\rightarrow$  Regression
- Cross-entropy  $\rightarrow$  Multi-class classification
- Cox model  $\rightarrow$  Survival models

# AdaBoost - Details

- Please notice that in AdaBoost, we are not minimizing the training error  $\epsilon_m$  given by the 0-1 loss but an upper bound, namely:  
$$\frac{1}{N} \sum_{i=1}^N I(F_M(x_i) \neq y_i) \leq \frac{1}{N} \sum_{i=1}^N \exp(-y_i h(x_i))$$
 where  $h(x_i) = \sum_m \alpha_m f_m(x_i)$
- Since the 0-1 loss is non-smooth and non-convex, it is common in statistical learning to use surrogate losses such as the exponential loss




Figure 4: Figure taken from [5].

- To increase generalization capability, we can also use Regularization techniques such as:
  - **Sampling** the training data (with or without replacement) using only a small part of the entire training set at each iteration  $\rightarrow$  it introduces randomness (models less similar/correlated), effective when data-set is very large
  - **Shrinking** the impact of the additional models. Simplest way is to introduce a hyper-parameter  $\gamma > 0$  in the update:  
$$h_m(x) = h_{m-1}(x) + \frac{1}{2}\gamma\alpha_m f(x; \beta_m) \rightarrow$$
to reduce model complexity and avoid overfitting
  - **Early stopping**, we choose the number of models  $f_m$  that corresponds to the minimum validation error

- XGBoost is a recent and very effective implementation of gradient boosting with some peculiarities:
  - It adds two regularization terms on  $f_m(x)$ : a L1-norm on the number of leaves, and a L2-norm on the leaf weights
  - it uses a second-order approximation of the cost function which gives a new impurity function
  - they also use shrinking (see slide before) and feature sub-sampling as in random forests
  - fast algorithm to find the best split (even for large data-sets)
  - can handle missing (or sparse) values
  - efficient implementation using a block structure and a cache-aware prefetching algorithm
  - Freely-available here:  
<https://xgboost.readthedocs.io/en/latest/build.html>

- LightGBM has been implemented to tackle data-sets with many features and large feature dimension
- Instead than scanning all observations for each feature to estimate the best split points, authors propose a technique to use only few observations (the ones with the greater gradients)
- Accuracy is preserved, with respect to other gradient boosting algorithms such as XGBoost but computational time is definitely lower
- Freely available here: <https://lightgbm.readthedocs.io/en/latest/Python-Intro.html>

## Other implementations

- **CatBoost**: efficient and fast implementation with support for categorical variables
- **pGBRT**: fast implementation especially used for ranking

**Take home message:** all of these implementations have many different hyper-parameters (for trees, regularization and optimization) that need to be set by the user.. not so simple ! It needs time and experience

**In scikit-learn**, you can find the implementation of AdaBoost and LightGBM (called Histogram-Based Gradient Boosting)

# Summary

## 1 Introduction

- Reminder about classification
- Constant piecewise model/classifier

## 2 Decision trees

- Cost functions
- Impurity function
- Stopping criteria and variations
- Model selection

## 3 Ensemble methods

- Bagging (Bootstrap aggregating)
- Random Forests
- Boosting
- Stacking



- Differently from Bagging and Boosting, Stacking considers **heterogeneous weak learners** (different models like KNN and SVM)
- Differently from Bagging and Boosting, Stacking aggregates the heterogeneous weak learners using an **external meta-model** (like a neural network) and not an deterministic strategy
- Depending on the problem, the user needs therefore to select the appropriate heterogeneous weak learners and the meta-model that combines their predictions

- ① Hastie, Tibshirani and Friedman (2009). The Element of Statistical Learning. Springer.
- ② Bishop (2006). Pattern Recognition and Machine Learning. Springer.
- ③ L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone (1984). Classification and regression trees. Wadsworth Statistics/Probability Series.
- ④ Breiman, L. (2001). Random forests, Machine Learning 45: 5–32.
- ⑤ R. Schapire and Y. Freund (2012). Boosting: Foundations and Algorithms. MIT Press.
- ⑥ Zhi-Hua Zhou (2012). Ensemble Methods: Foundations and Algorithms. Chapman & Hall/CRC