

Introduction to Deep Learning

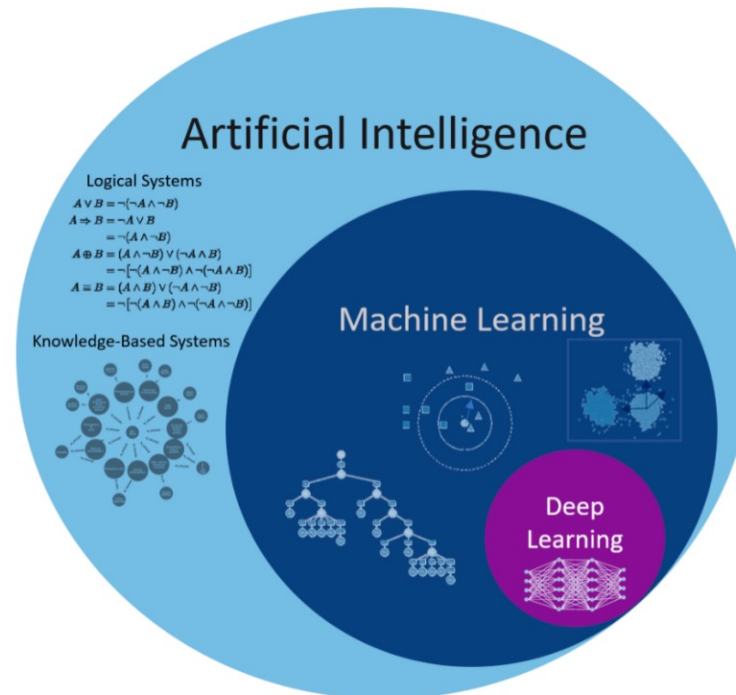
Multi-Layer Perceptrons (MLPs)

Loïc Le Folgoc

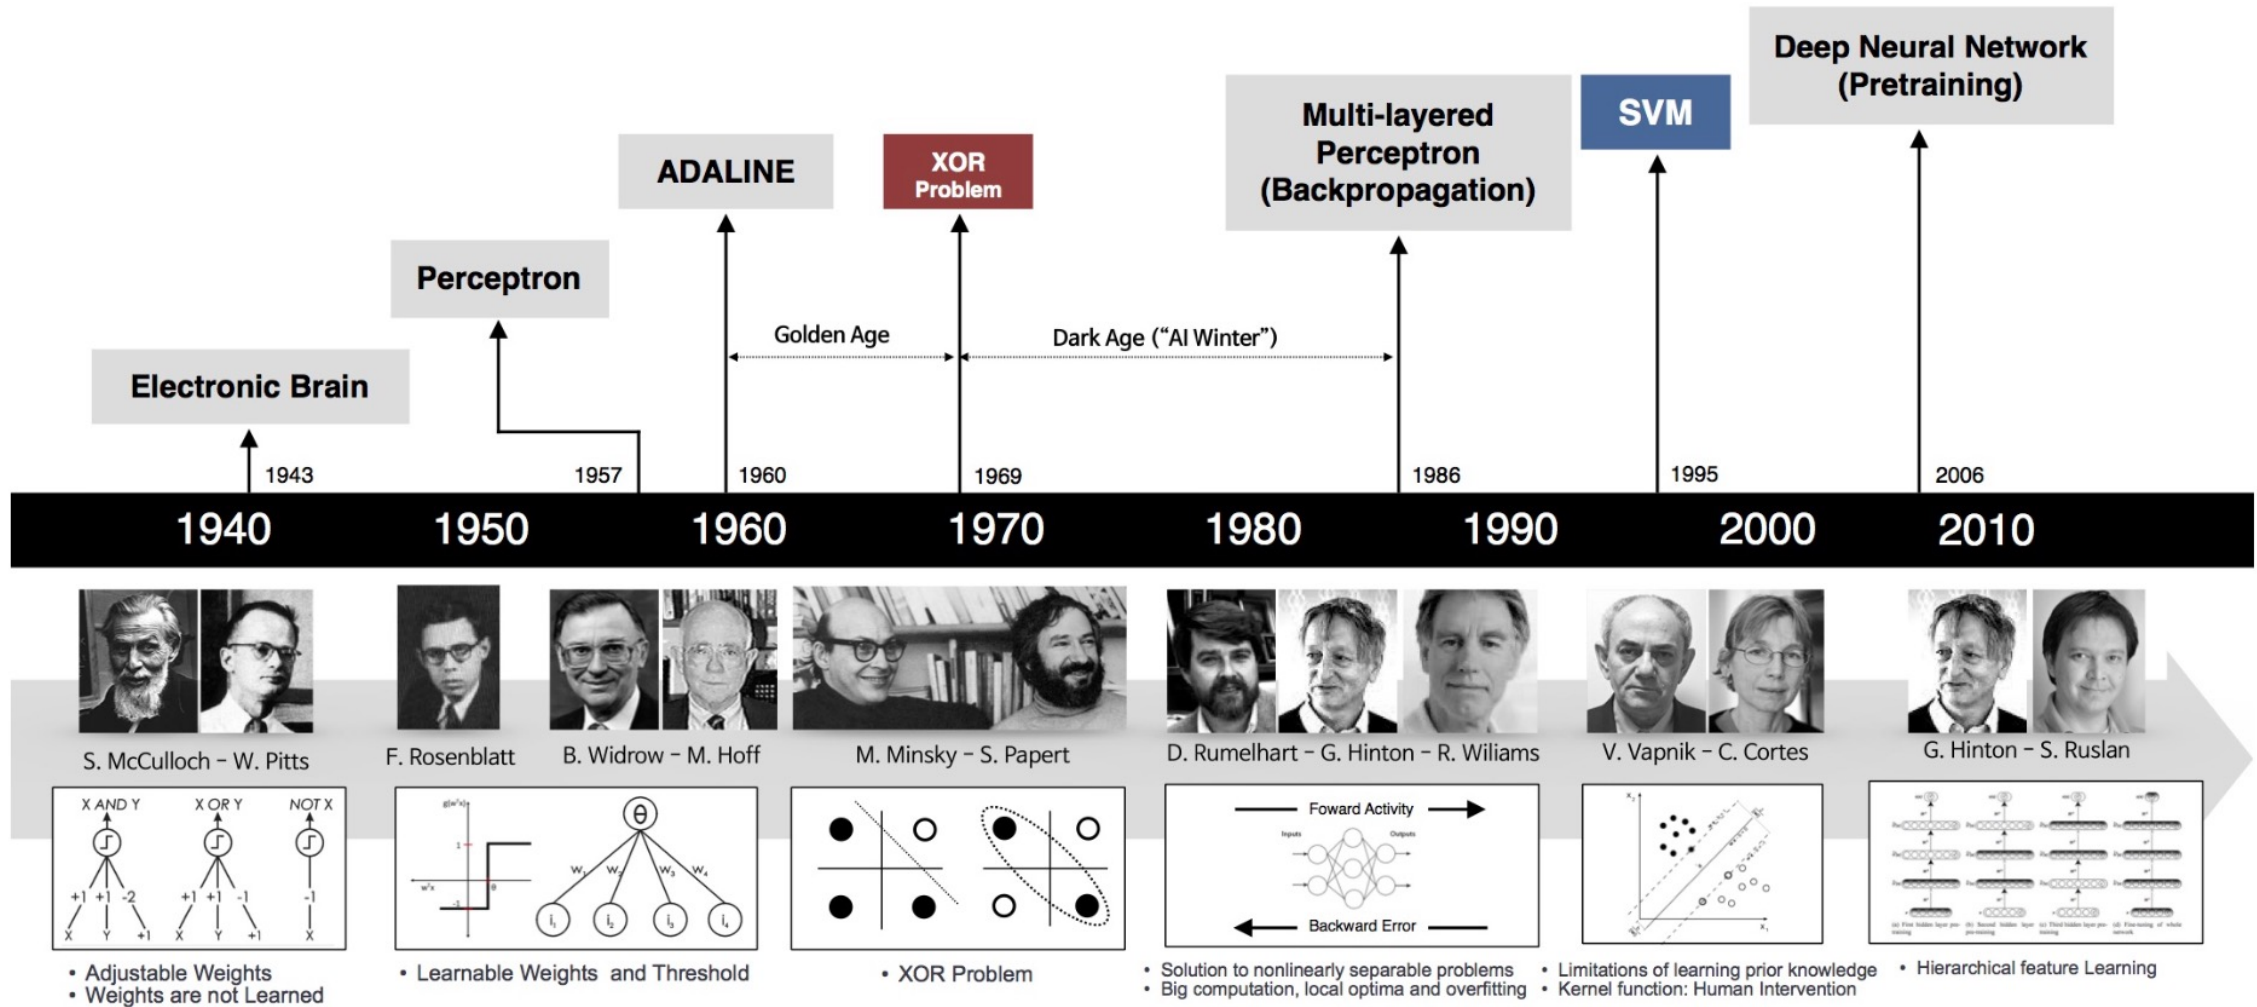
loic.lefolgoc@telecom-paris.fr

Introduction

- In this lesson, we will look at **Artificial Neural Networks** (ANNs), and in particular **Multi-Layer Perceptrons** (MLPs)
- It is a class of machine learning models with a wide range of applications, for supervised and unsupervised learning
- Recently, **deep learning** has had a huge impact on many domains, including image processing and computer vision



Evolution of AI and Deep Learning



Evolution of AI and Deep Learning


- Deep learning has been known by various names for a long time (“cybernetics”, artificial neural networks, etc.)
- Why has there been a recent explosion of deep learning?
 - Increase in number of large datasets (in terms of number of elements and dimensionality)
 - Increase in **computing power (GPUs)**
 - Auto-differentiation libraries

The Deep Learning “Computer Vision Recipe”



Big Data: ImageNet

+



Deep Convolutional Neural Network

+



Backprop on GPU

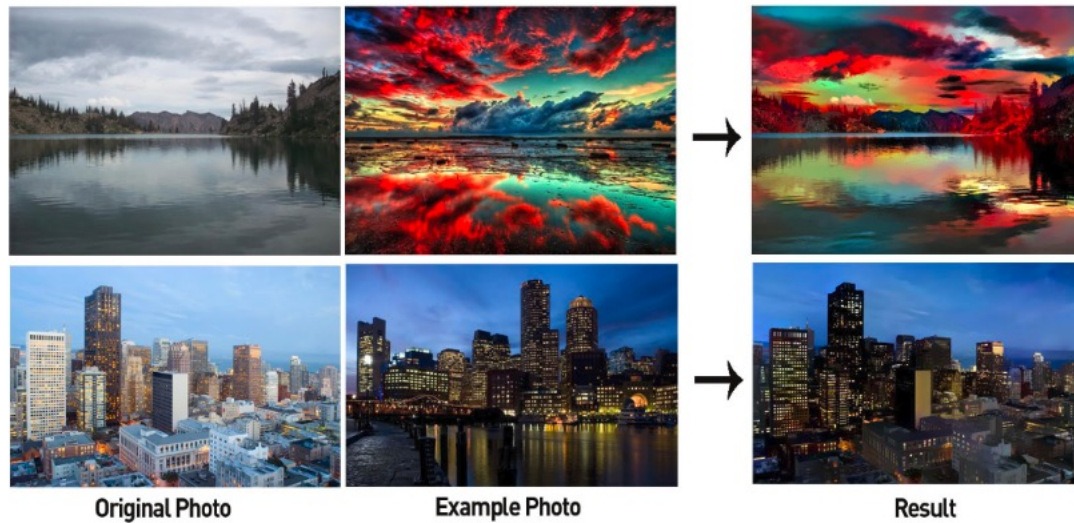
=



Learned Weights

Introduction

- Modern deep learning has applications in all areas of computer vision, and in a multitude of other domains
- Example: **style transfer**



Introduction

- Example: in 2016 **Google DeepMind's AlphaGo** beats Lee Sedol, one of the best player of Go in the world; in 2017 it beats Ke Jie, world #1



Artificial Neural networks / MLPs

Biological neuron

- The **brain** and its functioning **inspired early artificial intelligence**





Image source: wikipedia

- Neurons receive electric signals via *dendrites*, process the signal in the *soma* and transmit the output via *synapses*
- Neurons have an “all or none” principle: under a threshold, no electrical signal is emitted; above the threshold, the same one is always emitted

Biological neuron

- A **mathematical model of the neuron** was first proposed by McCulloch and Pitts*
- The model is quite simple (it is the simplest possible ANN):
 - A series of inputs
 - A list of weights
 - A threshold T (inspired by the “all or none” principle)




- The weights had to be set by hand; this is time consuming
- Rosenblatt** proposed an algorithm to optimize these weights, and named this model the **perceptron**

*A logical calculus of the ideas immanent in nervous activity, McCulloch, W., and Pitts, W., *The Bulletin of Mathematical Biophysics*, 1943

**The perceptron: A probabilistic model for information storage and organization in the brain, Rosenblatt, *Psychological Review*, 1958


First neural network: perceptron for binary classification

Hypothesis
$$h_{\theta}(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_i x_i \theta_i > 0 \\ 0 & \text{otherwise} \end{cases}$$




First neural network: perceptron for binary classification

Hypothesis
$$h_{\theta}(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_i x_i \theta_i > 0 \\ 0 & \text{otherwise} \end{cases}$$




First neural network: perceptron for binary classification

Hypothesis
$$h_{\theta}(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_i x_i \theta_i > 0 \\ 0 & \text{otherwise} \end{cases}$$



First neural network: perceptron for binary classification



Training a single-layer perceptron

Initialize the weights θ_i to 0 or a small random value

For each example in the training set $\mathbf{X} = \{\mathbf{x}^{(n)}\}_{n=1 \dots N}$ with labels $\{y^{(n)}\}_{n=1 \dots N}$:

Compute the prediction $\hat{y}^{(n)} \triangleq f(\boldsymbol{\theta}^T \mathbf{x}^{(n)})$

For all features $i = 1 \dots M$, update the weight:

$$\theta_i^{t+1} := \theta_i^t - \alpha \cdot (\hat{y}^{(n)} - y^{(n)}) x_i^{(n)}$$

Training a single-layer perceptron

Initialize the weights θ_i to 0 or a small random value

For each example in the training set $\mathbf{X} = \{\mathbf{x}^{(n)}\}_{n=1 \dots N}$ with labels $\{y^{(n)}\}_{n=1 \dots N}$:

Compute the prediction $\hat{y}^{(n)} \triangleq f(\boldsymbol{\theta}^T \mathbf{x}^{(n)})$

For all features $i = 1 \dots M$, update the weight:

$$\theta_i^{t+1} := \theta_i^t - \alpha \cdot (\hat{y}^{(n)} - y^{(n)}) x_i^{(n)}$$



learning rate

Training a single-layer perceptron

Initialize the weights θ_i to 0 or a small random value

For each example in the training set $\mathbf{X} = \{\mathbf{x}^{(n)}\}_{n=1\dots N}$ with labels $\{y^{(n)}\}_{n=1\dots N}$:

Compute the prediction $\hat{y}^{(n)} \triangleq f(\boldsymbol{\theta}^T \mathbf{x}^{(n)})$


For all features $i = 1 \dots M$, update the weight:

$$\theta_i^{t+1} := \theta_i^t - \alpha \cdot (\hat{y}^{(n)} - y^{(n)}) x_i^{(n)}$$

- Can be interpreted as **stochastic gradient descent** for a simple loss $\mathcal{L}(\boldsymbol{\theta}) \triangleq \sum_n (\hat{y}^{(n)} - y^{(n)}) \boldsymbol{\theta}^T \mathbf{x}^{(n)}$
- NN training will generalize this approach

The perceptron and the XOR function


- Unfortunately, a significant problem was found with the perceptron
- It is a linear classifier: Minsky and Papert* showed that the **XOR function** could not be approximated by a perceptron: the **XOR is not linearly separable**



- This resulted in a long period (about 10 or so years) when neural networks fell into disfavor; this was known as the “AI winter”

Multi-Layer Perceptron (MLP)

- However, a solution was found by introducing several layers (“hidden layers”) → **MLP**
 - Rumelhart et al.* showed that this could be trained to learn the XOR function

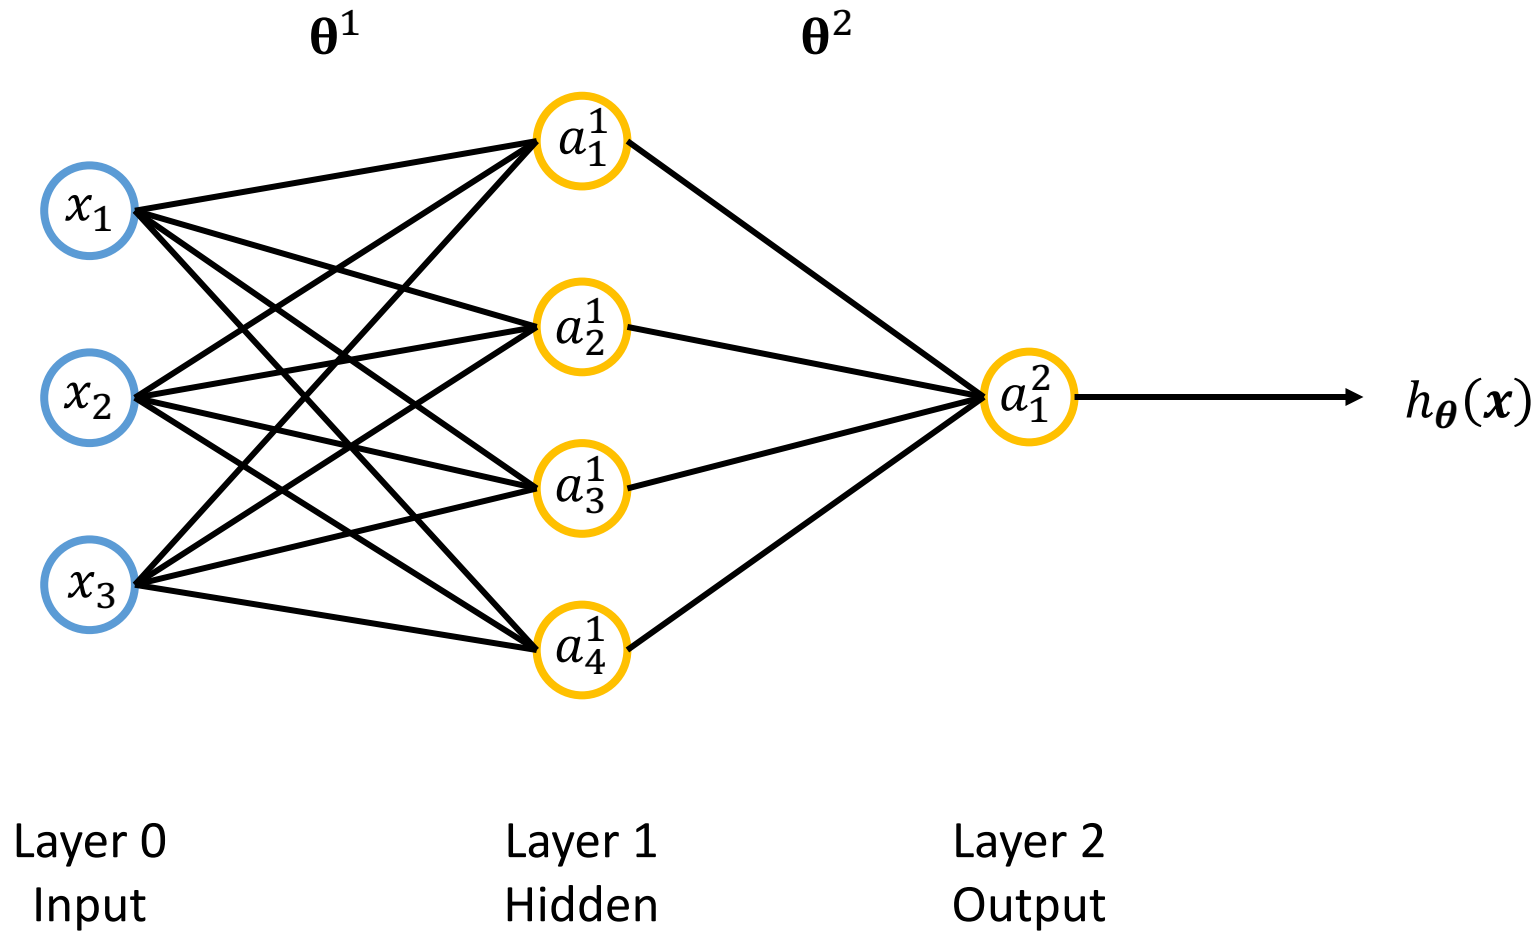


- In fact, the **universal approximation theorem**** guarantees that any continuous function may be approximated with arbitrary precision by a large enough MLP with 1 hidden layer (although the network may be unfeasibly large)

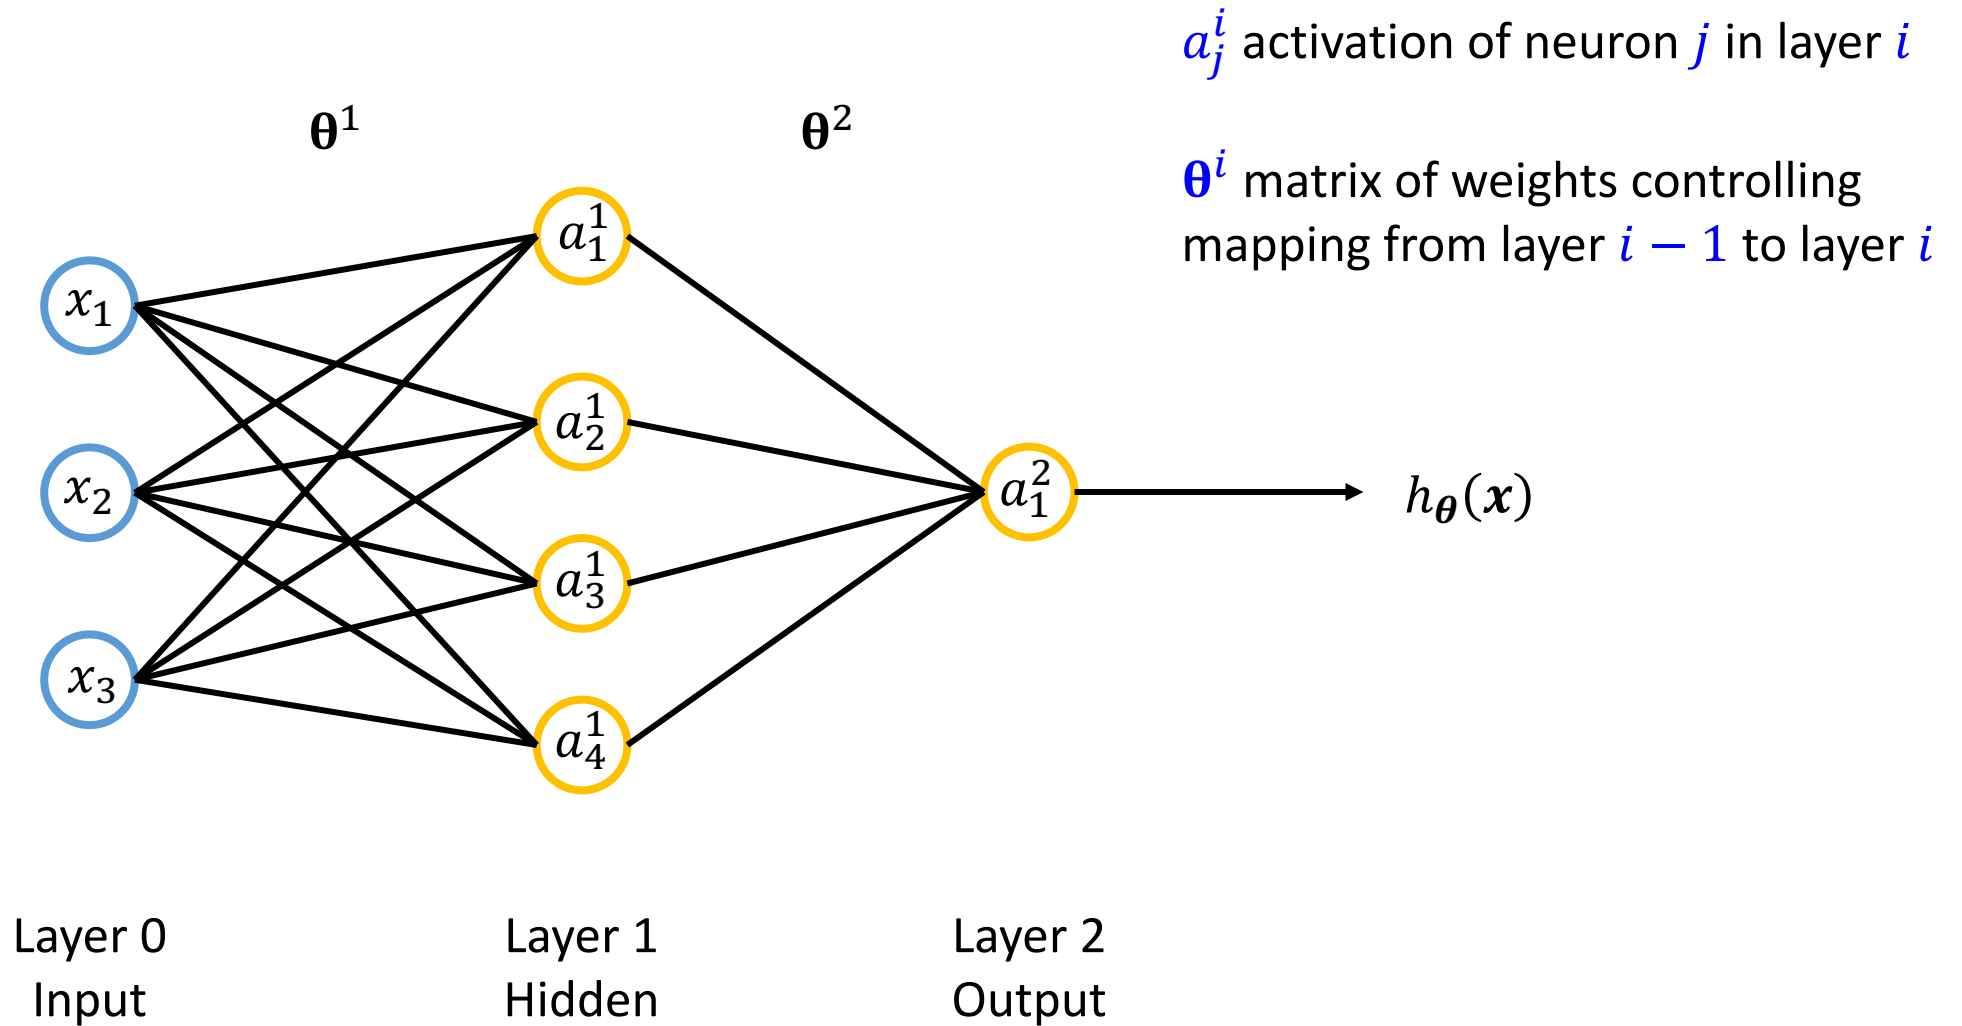
***Learning internal representations by error propagation**, Rumelhart, D., Hinton, G., and Williams, R., *California University San Diego LA Jolla Inst. For Cognitive Science*, 1985

****Approximations by superposition of sigmoidal functions**, Cybenko, G., *Mathematics of Control, Signals and Systems*, 1989

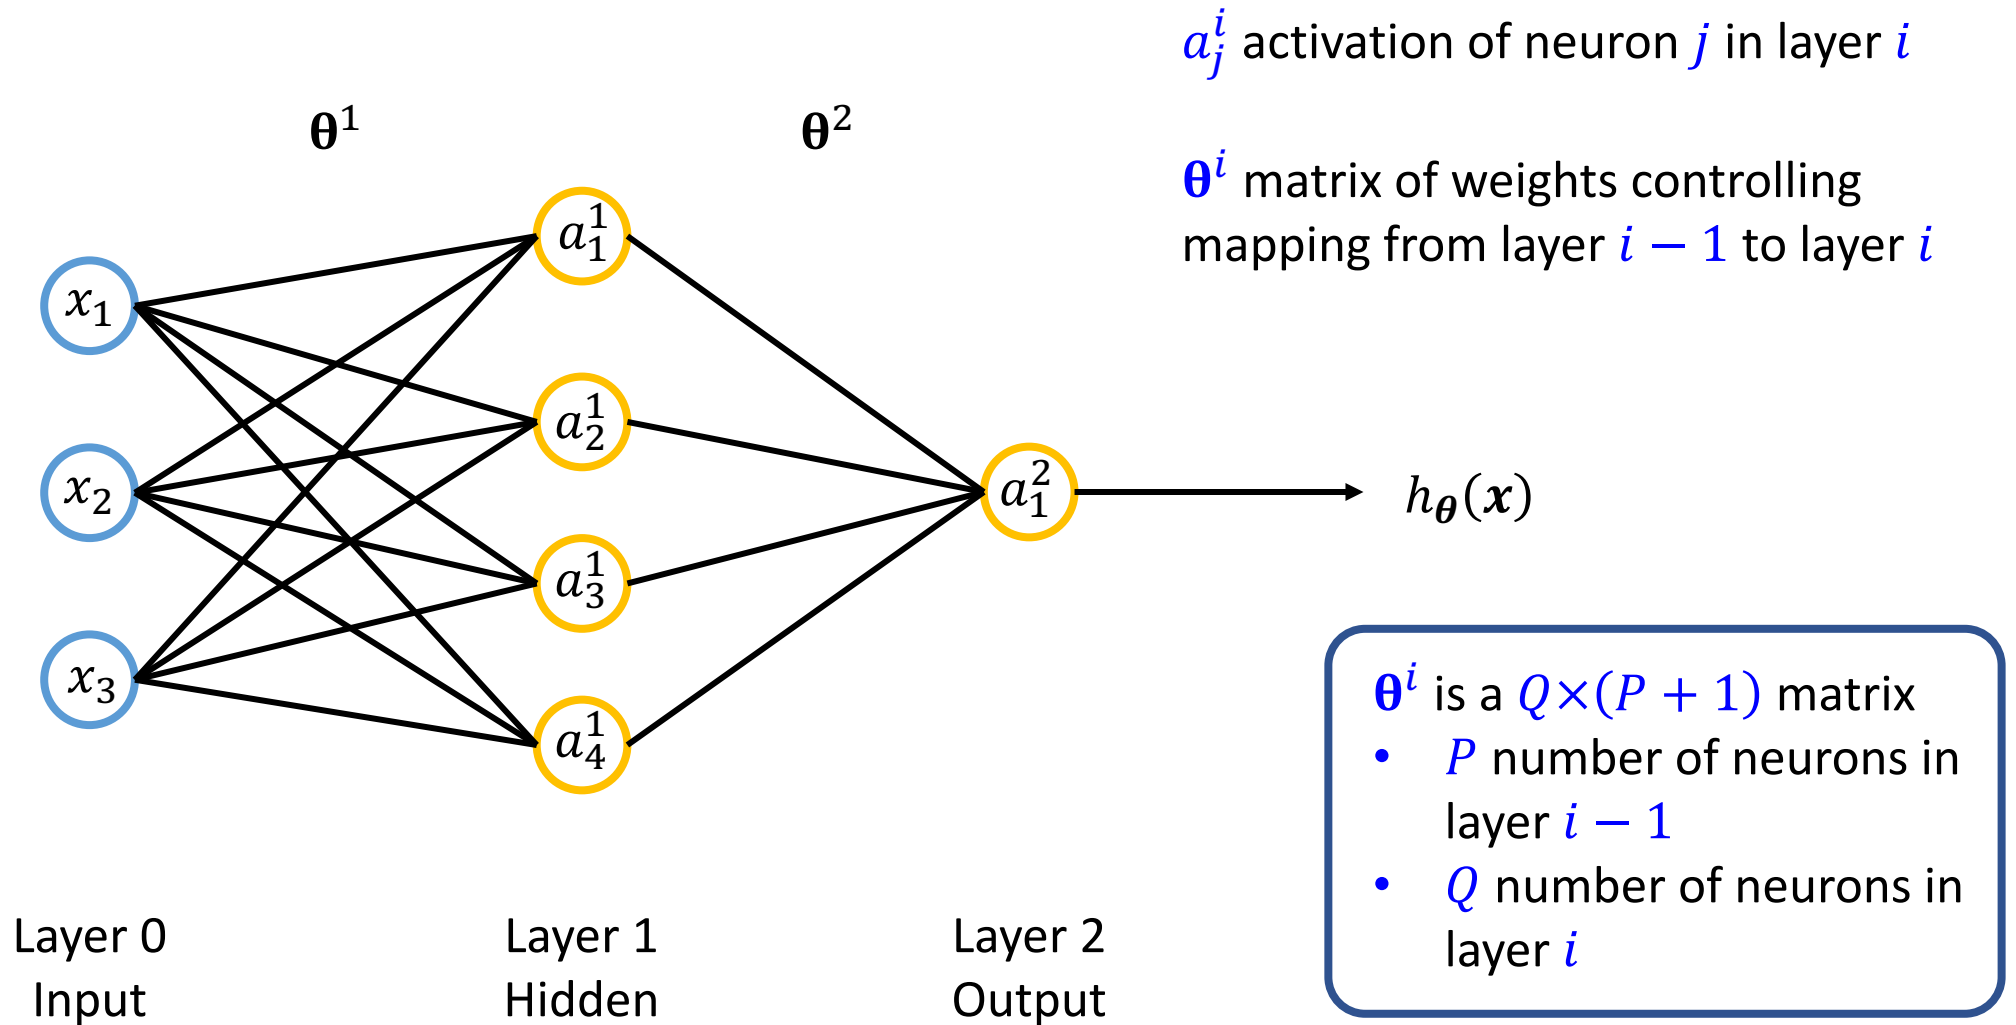
Multi-Layer Perceptron




Multi-Layer Perceptron



Multi-Layer Perceptron



Multi-Layer Perceptron



$$\mathbf{a}^1 := f(\boldsymbol{\theta}^1 \mathbf{x})$$

$$\mathbf{a}^2 := f(\boldsymbol{\theta}^2 \mathbf{a}^1) = f(\boldsymbol{\theta}^2 f(\boldsymbol{\theta}^1 \mathbf{x}))$$

$$a_1^1 := f(\theta_{0,0}^1 + \theta_{0,1}^1 x_1 + \theta_{0,2}^1 x_2 + \theta_{0,3}^1 x_3)$$

$$a_2^1 := f(\theta_{1,0}^1 + \theta_{1,1}^1 x_1 + \theta_{1,2}^1 x_2 + \theta_{1,3}^1 x_3)$$

$$a_3^1 := f(\theta_{2,0}^1 + \theta_{2,1}^1 x_1 + \theta_{2,2}^1 x_2 + \theta_{2,3}^1 x_3)$$

$$a_4^1 := f(\theta_{3,0}^1 + \theta_{3,1}^1 x_1 + \theta_{3,2}^1 x_2 + \theta_{3,3}^1 x_3)$$


$$h_{\theta}(x) := a_1^2 := f(\theta_{0,0}^2 + \theta_{0,1}^2 a_1^1 + \theta_{0,2}^2 a_2^1 + \theta_{0,3}^2 a_3^1 + \theta_{0,4}^2 a_4^1)$$

Multiclass classification: NN with multiple outputs

- So far we have considered NNs with a single output; this works for:
 - Scalar regression $y \in \mathbb{R}$
 - Binary classification $y \in \{0,1\}$
- For K classes we can use a vector $\mathbf{y} = (y_1, \dots, y_K)$ such that
$$y_k := \begin{cases} 1 & \text{if } k \text{ is the index of the true class} \\ 0 & \text{otherwise} \end{cases}$$
- This is called **one-hot encoding** (only one element is different from 0)
- NN outputs can represent class probabilities

Multiclass classification: NN with multiple outputs

With $K = 3$



Softmax layer

The **softmax function** transforms K real values to a K -vector of probabilities summing to 1

$$\hat{y}_k \triangleq \frac{e^{a_k}}{\sum_l e^{a_l}}$$

The outputs \hat{y}_k are such that

- $\hat{y}_k \geq 0$
- $\sum_k \hat{y}_k = 1$

This is an **extension of the sigmoid** function to K classes

Loss function for K-way classification

How to train? Minimize **cross-entropy loss**, also called log-loss

- In general, cross-entropy $H(p, q)$ of probability distributions p and q

$$H(p, q) \triangleq - \sum_{k=1}^K p_k \log q_k$$

- We define the loss for one data point as $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) \triangleq H(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{k=1}^K y_k \log \hat{y}_k$
 - where the prediction $\hat{\mathbf{y}}$ depends on the parameters of the model

Probabilistic interpretation:

- This is equivalent to modeling the class by a **categorical distribution** with probabilities \hat{y}_k :

$$\mathcal{C}(\mathbf{y}|\hat{\mathbf{y}}) = \prod_{k=1}^K \hat{y}_k^{y_k}$$

- then taking the **negative log-likelihood** $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) \triangleq - \log \mathcal{C}(\mathbf{y}|\hat{\mathbf{y}})$ as loss function

Loss function for regression

Minimize **mean squared error**, also called sum-of-squared-differences


- We define the loss for one data point as $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) \triangleq (\hat{\mathbf{y}} - \mathbf{y})^2$
 - where the prediction $\hat{\mathbf{y}}$ depends on the parameters of the model
- The total loss over the training dataset is a sum/mean of individual losses over each data point

Activation functions f for hidden layers

Heaviside function

- Not differentiable

$$f(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{otherwise} \end{cases}$$




Activation functions f for hidden layers

Sigmoid function

- Squashes numbers to range (0,1)
- Neurons can saturate:
values close to 0 or 1
- Saturated neurons “kill” gradients
= “vanishing gradients” issue
- Sigmoid outputs are not 0 centered
 - Ok for last layer (for binary classification)
 - Not for intermediate layers
- Exp is expensive to compute

$$f(z) = \frac{1}{1 + e^{-z}}$$

$$f'(z) = f(z)(1 - f(z))$$




Activation functions f for hidden layers

tanh function

- Rescaled and centered variant of sigmoid
- Squashes numbers to range $(-1,1)$
- tanh outputs are 0 centered
 - Good for intermediate layers
- Neurons can saturate:
values close to -1 or 1
- Saturated neurons “kill” gradients
= “vanishing gradients” issue
- tanh is expensive to compute

$$f(z) = \tanh z = \frac{2}{1 + e^{-2z}} - 1$$

$$f'(z) = 1 - f(z)^2$$




Activation functions f for hidden layers

ReLU (Rectified Linear Unit) function

- Piecewise linear
- Does not saturate (positive values)
- Computationally cheap
- Converges faster than sigmoid/tanh in practice ($\approx 6\times$)
- Not zero-centered output
- “Dead neurons”: no gradient to drive parameter update when input is below 0

$$f(z) = \max(0, z)$$

$$f'(z) = \begin{cases} 0 & \text{for } z < 0 \\ 1 & \text{for } z \geq 0 \end{cases}$$




Activation functions f for hidden layers

Leaky ReLU function

- Piecewise linear
- Does not saturate
- Computationally cheap
- Converges faster than sigmoid/tanh in practice ($\approx 6\times$)
- Will not die (no “dead neurons”)
- Closer to 0-mean outputs
- PReLU (Parametric ReLU) variant allows to learn slope α based on data

$$f(z) = \max(\alpha z, z)$$

$$f'(z) = \begin{cases} \alpha & \text{for } z < 0 \\ 1 & \text{for } z \geq 0 \end{cases}$$




Activation functions f for hidden layers

ELU (Exponential Linear Unit) function

- Does not saturate (positive values)
- Computationally cheap in positive half
- Will not die (no “dead neurons”, although it does saturate in the negative half)
- Closer to 0-mean outputs
- Computationally expensive in negative half

$$f(z) = \begin{cases} z & \text{if } z > 0 \\ \alpha(\exp(z) - 1) & \text{if } z \leq 0 \end{cases}$$

$$f'(z) = \begin{cases} 1 & \text{if } z > 0 \\ f(z) + \alpha & \text{if } z \leq 0 \end{cases}$$




Activation functions f for hidden layers

Softplus function

- Range $(0, +\infty)$
- Useful as last layer when predicting a positive scalar (standard deviation, variance)
- Does not saturate (positive values)
- Will not die (no “dead neurons”, although it does saturate in the negative half)
- Not 0-centered outputs
- Computationally expensive

$$f(z) = \ln(1 + e^z)$$

$$f'(z) = \frac{1}{1 + e^{-z}}$$



Training Neural networks: backpropagation and computational graphs

Training neural networks

Define a loss function:

$$\mathcal{L}(\boldsymbol{\theta}) \triangleq \frac{1}{N} \sum_{n=1}^N \mathcal{L}(\mathbf{y}^{(n)}, \hat{\mathbf{y}}(\mathbf{x}^{(n)}, \boldsymbol{\theta}))$$

Optimize the loss, based on gradient descent:

repeat until convergence:

$$\boldsymbol{\theta}^{t+1} := \boldsymbol{\theta}^t - \alpha \cdot \nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}^t)$$

Training neural networks

Define a loss function:

$$\mathcal{L}(\boldsymbol{\theta}) \triangleq \frac{1}{N} \sum_{n=1}^N \mathcal{L}(\mathbf{y}^{(n)}, \hat{\mathbf{y}}(\mathbf{x}^{(n)}, \boldsymbol{\theta}))$$

Optimize the loss, based on gradient descent:

repeat until convergence:

$$\boldsymbol{\theta}^{t+1} := \boldsymbol{\theta}^t - \alpha \cdot \nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}^t)$$

How to compute the $\frac{\partial \mathcal{L}}{\partial \theta_i}(\boldsymbol{\theta}^t)$ for all parameters i ?

From the chain rule to forward and backward passes

Consider $\theta \mapsto f(\theta)$, $y \mapsto g(y)$, $z \mapsto h(z)$ and a function $\mathcal{L}(\theta) \triangleq h(g(f(\theta)))$.

Suppose we want to compute $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*)$

By the chain rule,

$$\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h}{\partial z}(g(f(\theta^*))) \cdot \frac{\partial g}{\partial y}(f(\theta^*)) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$$

- We need to compute, in order, from θ^* : $y^* := f(\theta^*)$, $z^* := g(y^*) \rightarrow$ **“forward pass”**
- **“Backward pass”**: we can then compute, in order:
 1. $\frac{\partial h}{\partial z}(z^*)$
 2. $\frac{\partial h \circ g}{\partial y}(y^*) = \frac{\partial h}{\partial z}(z^*) \cdot \frac{\partial g}{\partial y}(y^*)$
 3. $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h \circ g}{\partial y}(y^*) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$

From the chain rule to forward and backward passes

Consider $\theta \mapsto f(\theta)$, $y \mapsto g(y)$, $z \mapsto h(z)$ and a function $\mathcal{L}(\theta) \triangleq h(g(f(\theta)))$.

Suppose we want to compute $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*)$

By abuse of notation,
we also write:

- $\mathcal{L}(z) \triangleq h(z)$
- $\mathcal{L}(y) \triangleq h(g(y))$

By the chain rule,

$$\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h}{\partial z}(g(f(\theta^*))) \cdot \frac{\partial g}{\partial y}(f(\theta^*)) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$$

- We need to compute, in order, from θ^* : $y^* := f(\theta^*)$, $z^* := g(y^*) \rightarrow$ **“forward pass”**
- **“Backward pass”**: we can then compute, in order:
 1. $\frac{\partial h}{\partial z}(z^*)$
 2. $\frac{\partial h \circ g}{\partial y}(y^*) = \frac{\partial h}{\partial z}(z^*) \cdot \frac{\partial g}{\partial y}(y^*)$
 3. $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h \circ g}{\partial y}(y^*) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$

From the chain rule to forward and backward passes

Consider $\theta \mapsto f(\theta)$, $y \mapsto g(y)$, $z \mapsto h(z)$ and a function $\mathcal{L}(\theta) \triangleq h(g(f(\theta)))$.

Suppose we want to compute $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*)$

By abuse of notation,
we also write:

- $\mathcal{L}(z) \triangleq h(z)$
- $\mathcal{L}(y) \triangleq h(g(y))$

By the chain rule,

$$\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h}{\partial z}(g(f(\theta^*))) \cdot \frac{\partial g}{\partial y}(f(\theta^*)) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$$

- We need to compute, in order, from θ^* : $y^* := f(\theta^*)$, $z^* := g(y^*) \rightarrow$ **“forward pass”**
- **“Backward pass”**: we can then compute, in order:

1. $\frac{\partial h}{\partial z}(z^*) \rightarrow$ we also note this quantity $\frac{\partial \mathcal{L}}{\partial z}$
2. $\frac{\partial h \circ g}{\partial y}(y^*) = \frac{\partial h}{\partial z}(z^*) \cdot \frac{\partial g}{\partial y}(y^*) \rightarrow$ we also note this quantity $\frac{\partial \mathcal{L}}{\partial y} = \frac{\partial \mathcal{L}}{\partial z} \frac{\partial z}{\partial y}$
3. $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h \circ g}{\partial y}(y^*) \cdot \frac{\partial f}{\partial \theta}(\theta^*) \rightarrow \frac{\partial \mathcal{L}}{\partial \theta} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial \theta}$

From the chain rule to forward and backward passes

Consider $\theta \mapsto f(\theta)$, $y \mapsto g(y)$, $z \mapsto h(z)$ and a function $\mathcal{L}(\theta) \triangleq h(g(f(\theta)))$.

Suppose we want to compute $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*)$

By abuse of notation,
we also write:

- $\mathcal{L}(z) \triangleq h(z)$
- $\mathcal{L}(y) \triangleq h(g(y))$

By the chain rule,

$$\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h}{\partial z}(g(f(\theta^*))) \cdot \frac{\partial g}{\partial y}(f(\theta^*)) \cdot \frac{\partial f}{\partial \theta}(\theta^*)$$


- We need to compute, in order, from θ^* : $y^* := f(\theta^*)$, $z^* := g(y^*) \rightarrow$ **“forward pass”**

- **“Backward pass”**: we can then compute, in order:

1. $\frac{\partial h}{\partial z}(z^*) \rightarrow$ we also note this quantity $\frac{\partial \mathcal{L}}{\partial z}$
2. $\frac{\partial h \circ g}{\partial y}(y^*) = \frac{\partial h}{\partial z}(z^*) \cdot \frac{\partial g}{\partial y}(y^*) \rightarrow$ we also note this quantity $\frac{\partial \mathcal{L}}{\partial y}$
3. $\frac{\partial \mathcal{L}}{\partial \theta}(\theta^*) = \frac{\partial h \circ g}{\partial y}(y^*) \cdot \frac{\partial f}{\partial \theta}(\theta^*) \rightarrow \frac{\partial \mathcal{L}}{\partial \theta}$

Compute the derivative of the output
w.r.t. all intermediate variables,
starting from the “deepest” variable

Computational graphs, an example



$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1x_2 - x_1x_3 + x_2^2 - x_2x_3$$


$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example

forward pass



$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1x_2 - x_1x_3 + x_2^2 - x_2x_3$$


$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example

forward pass



$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1x_2 - x_1x_3 + x_2^2 - x_2x_3$$


$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example

forward pass



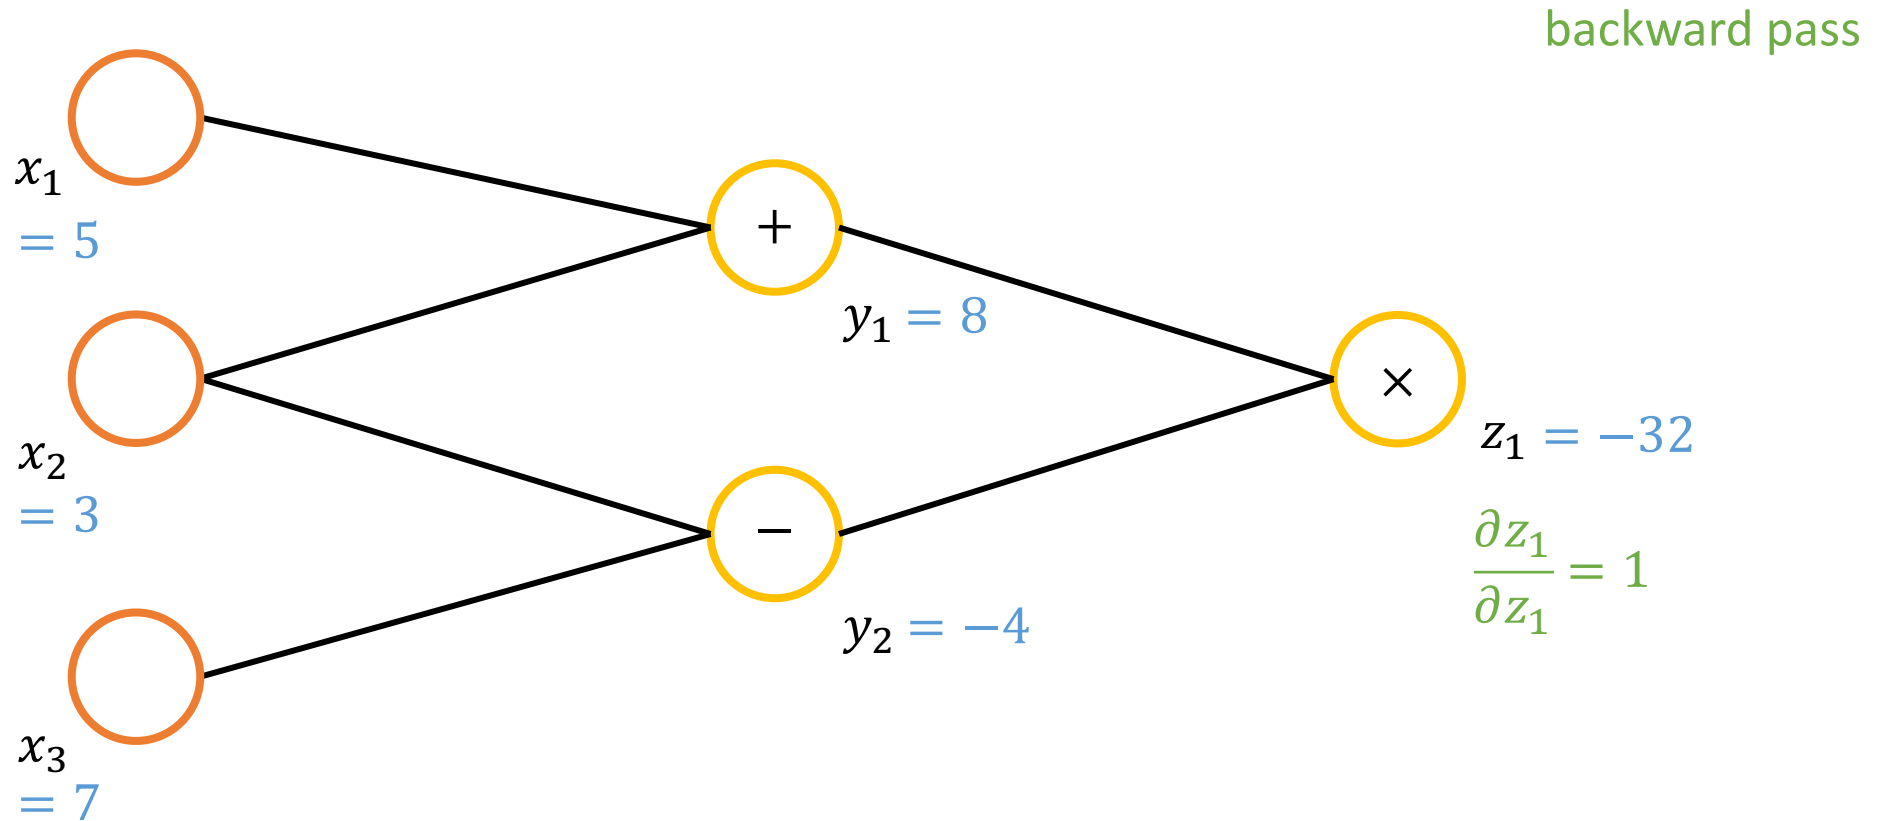
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



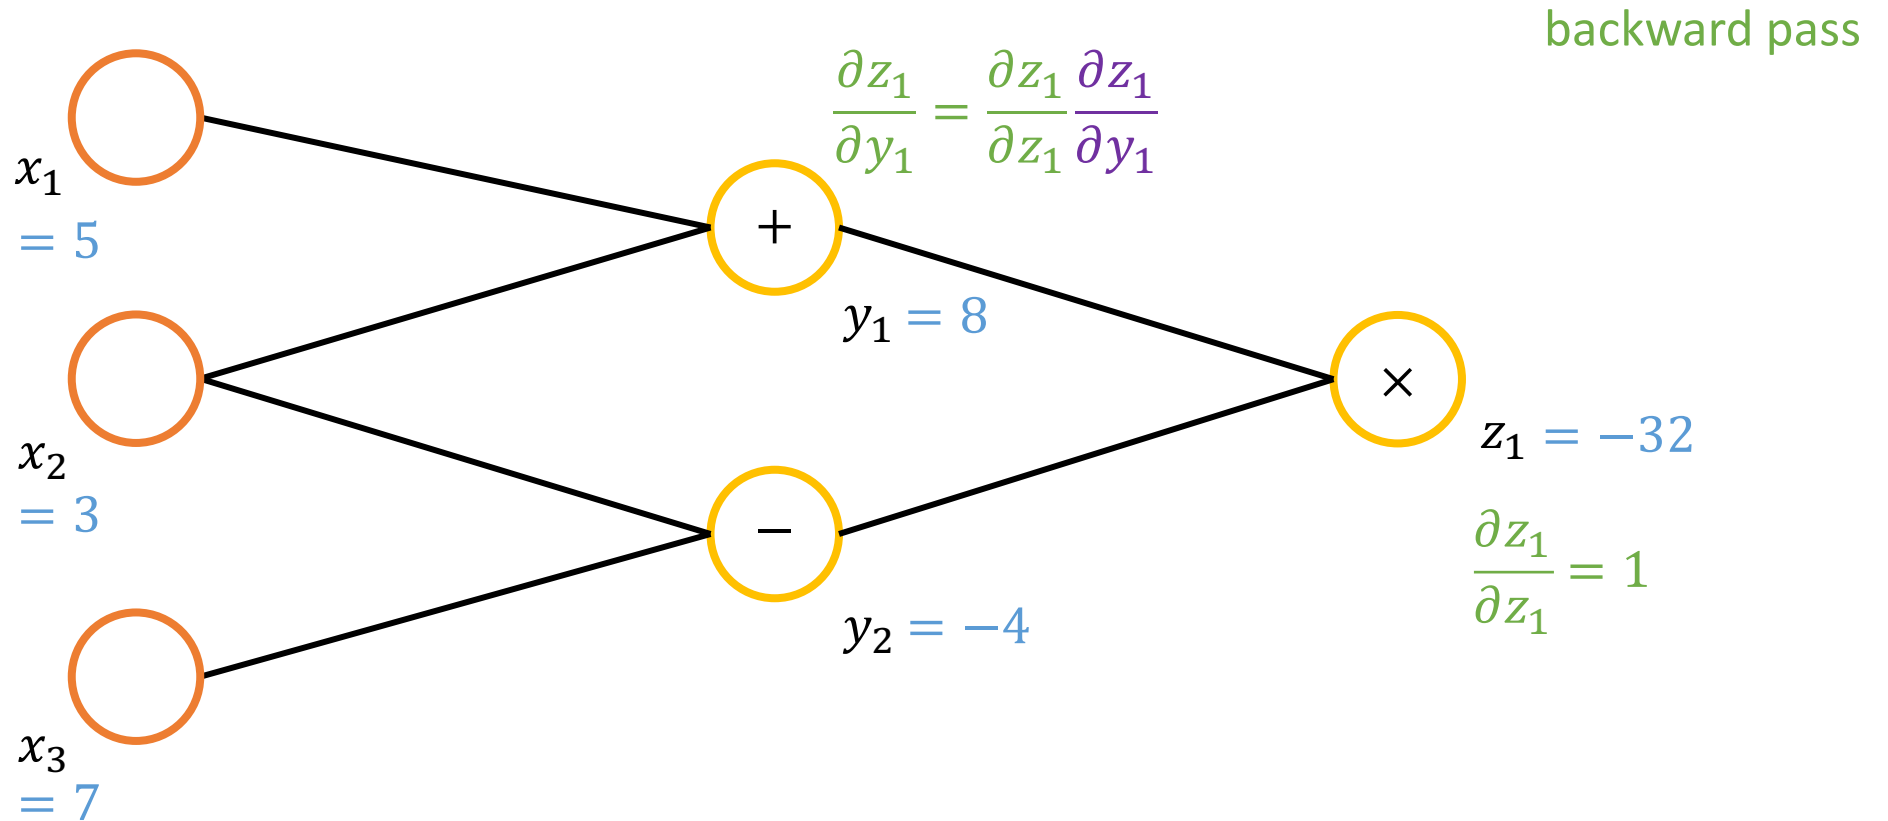
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



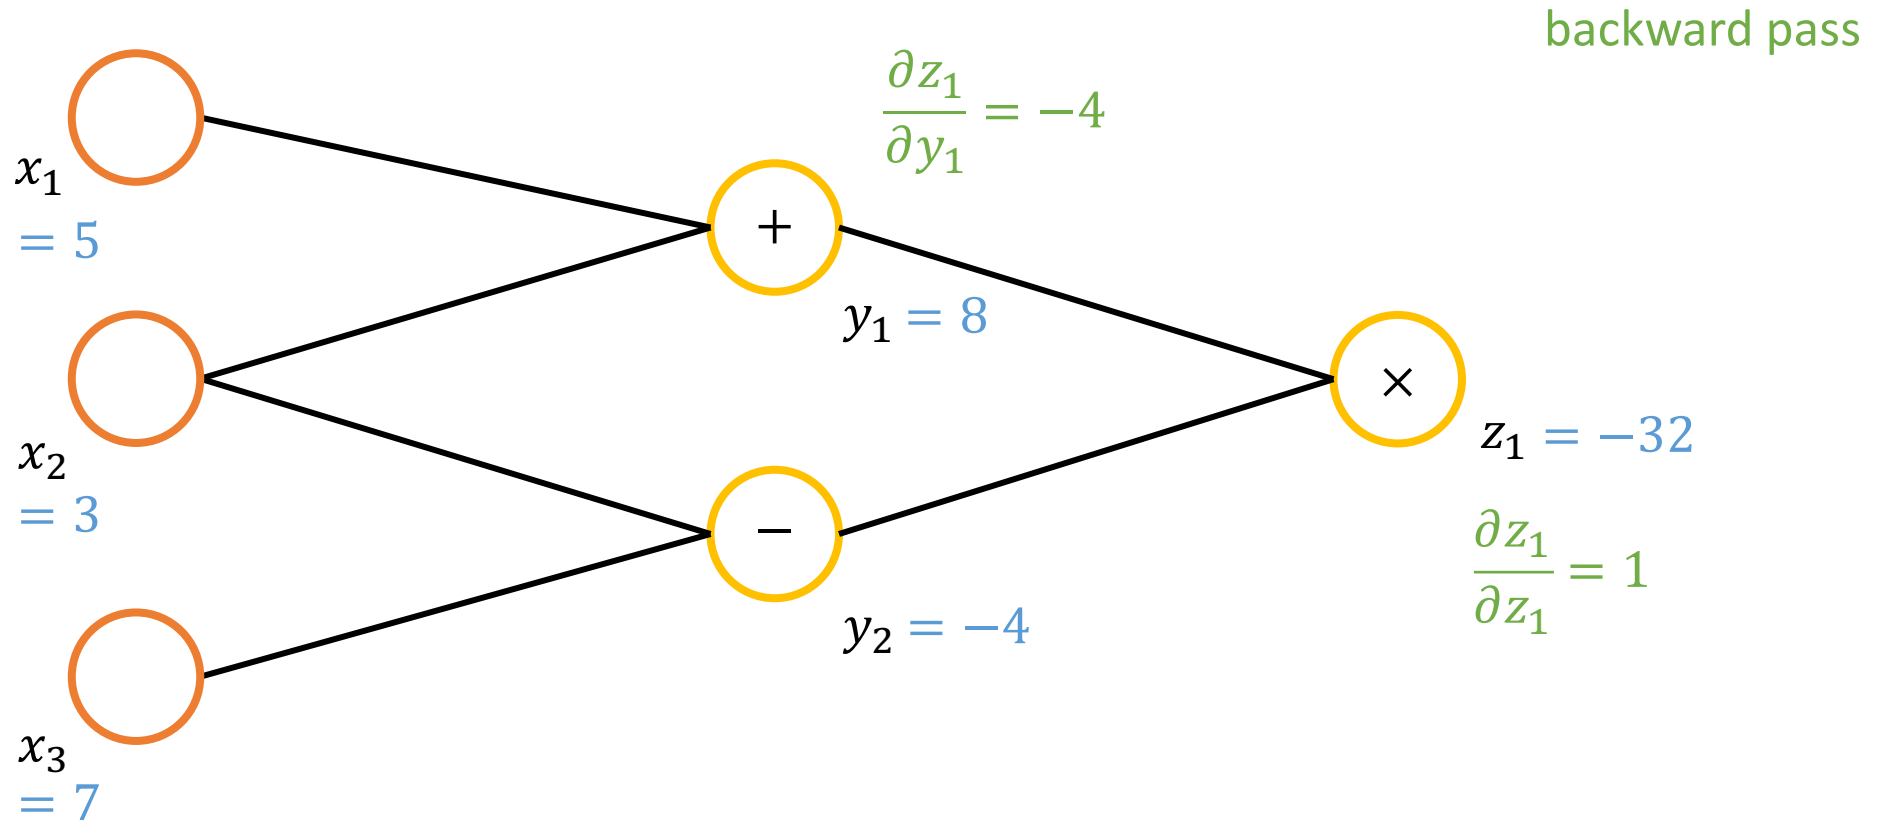
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



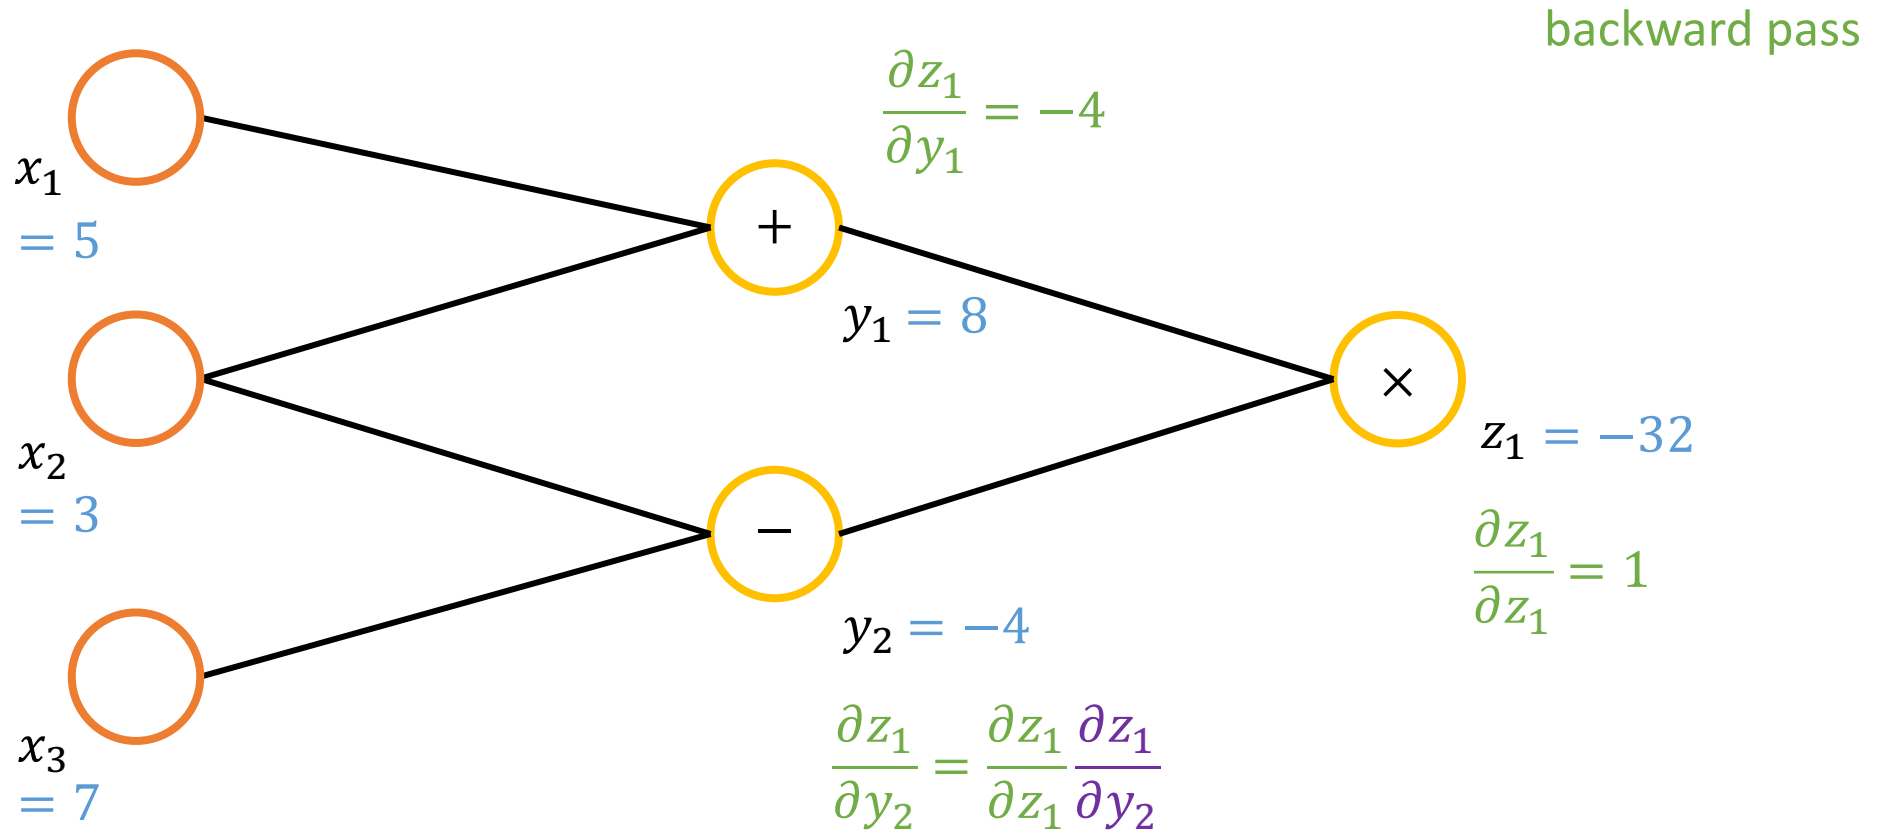
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



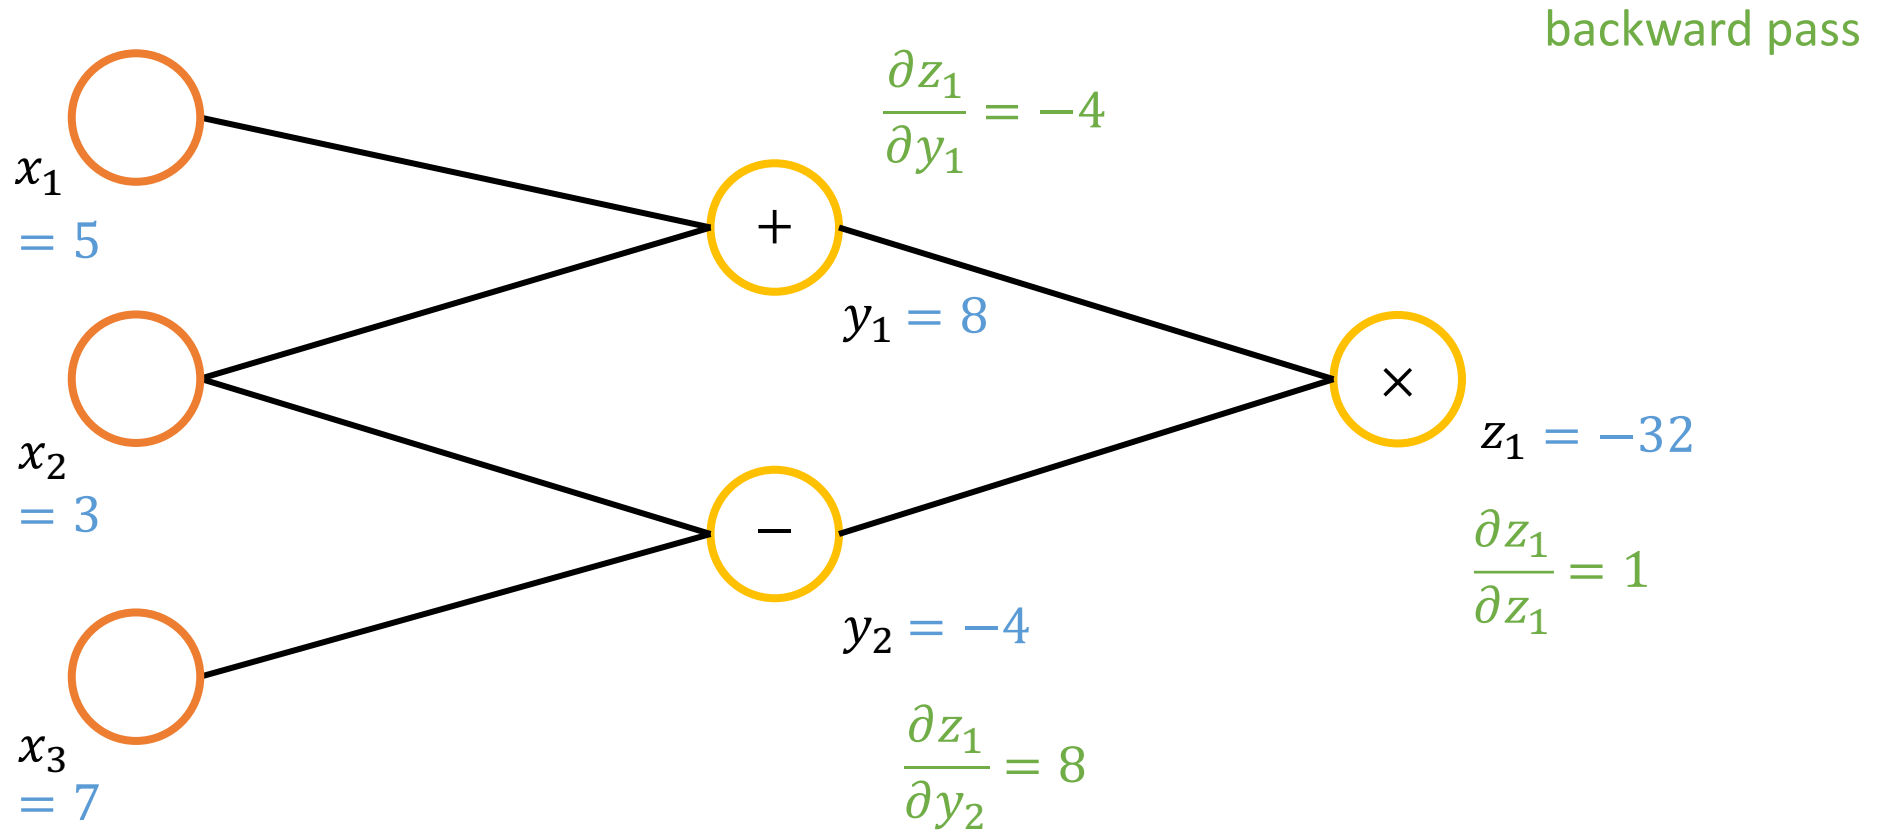
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



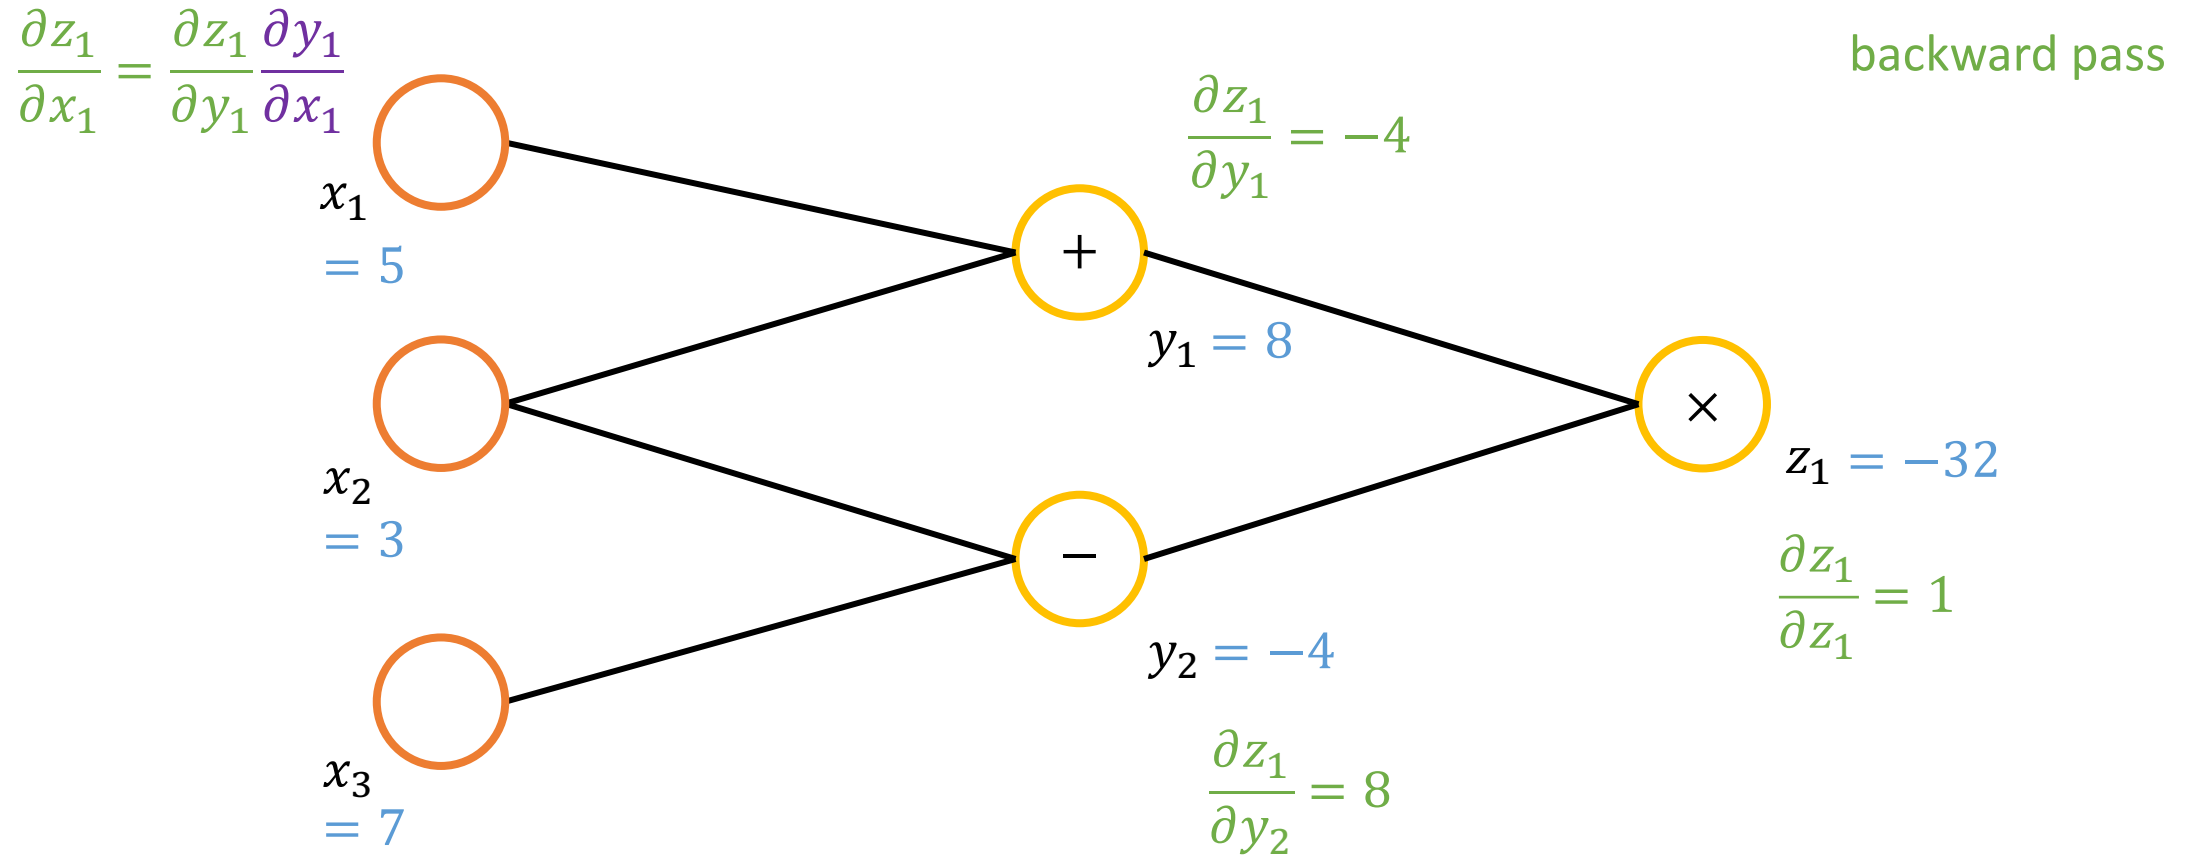
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example




$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



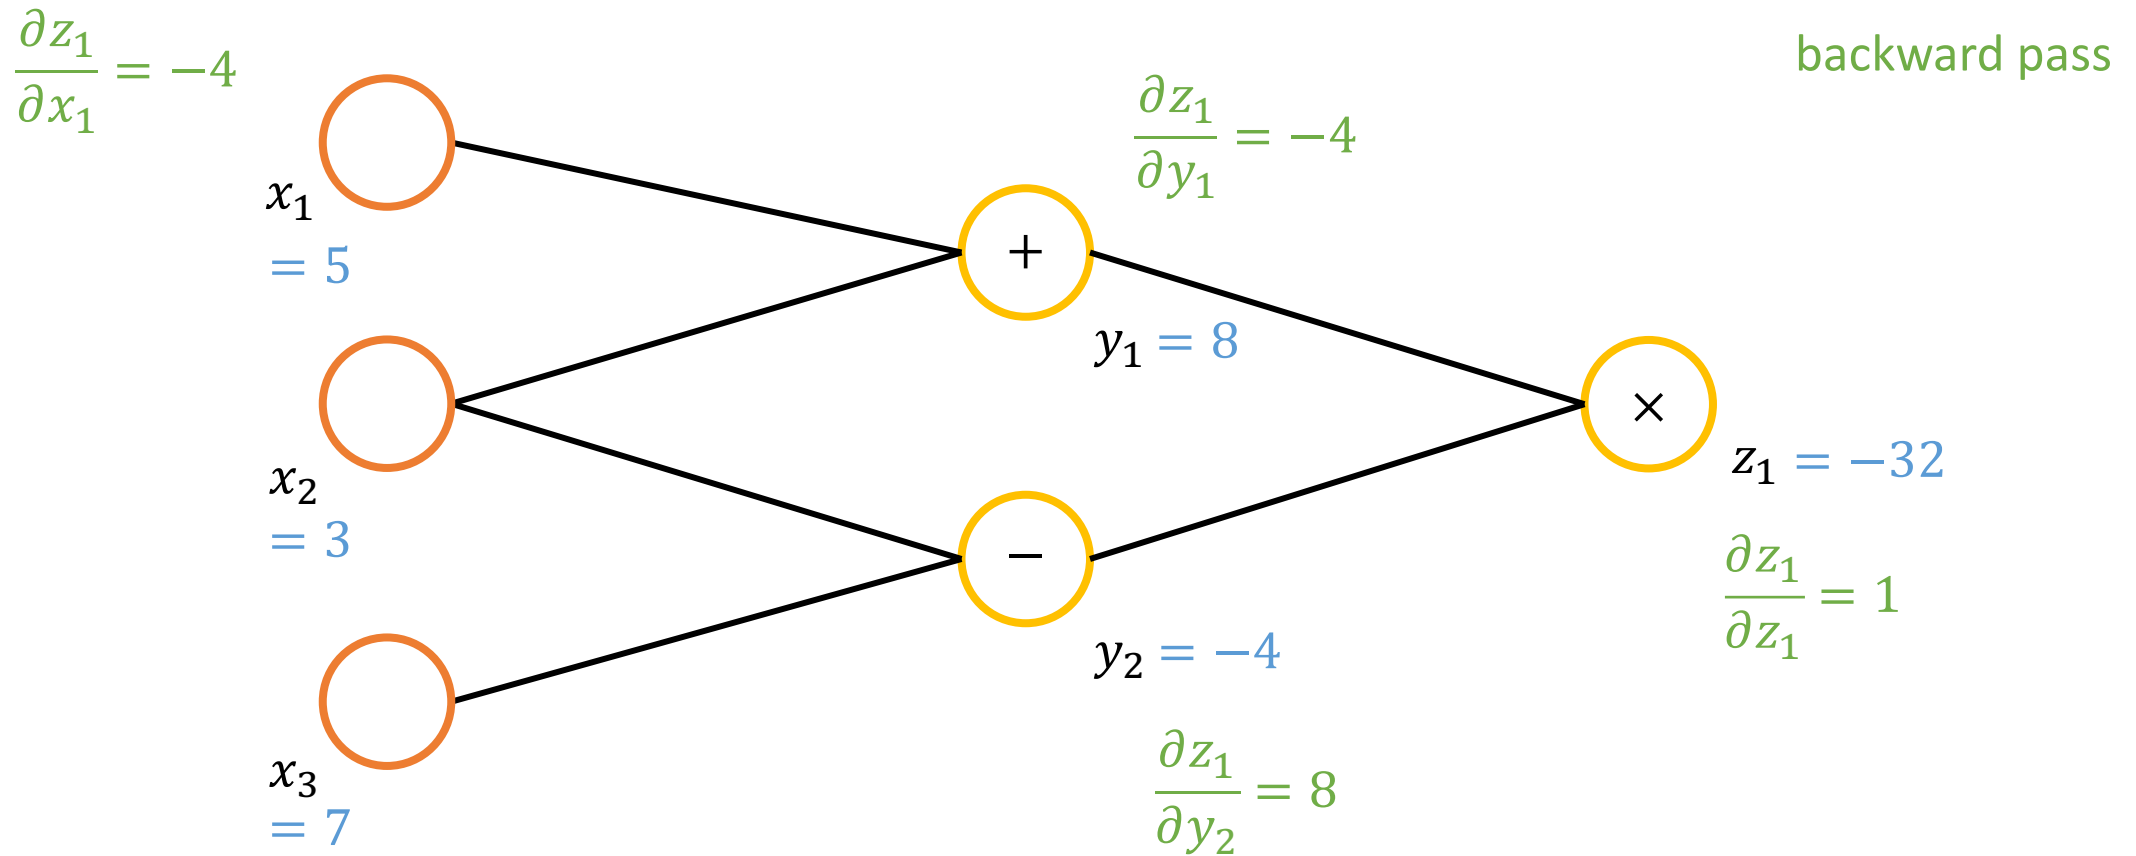
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



$$\frac{\partial z_1}{\partial x_3} = \frac{\partial z_1}{\partial y_2} \frac{\partial y_2}{\partial x_3}$$


$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



$$\frac{\partial z_1}{\partial x_3} = -8$$

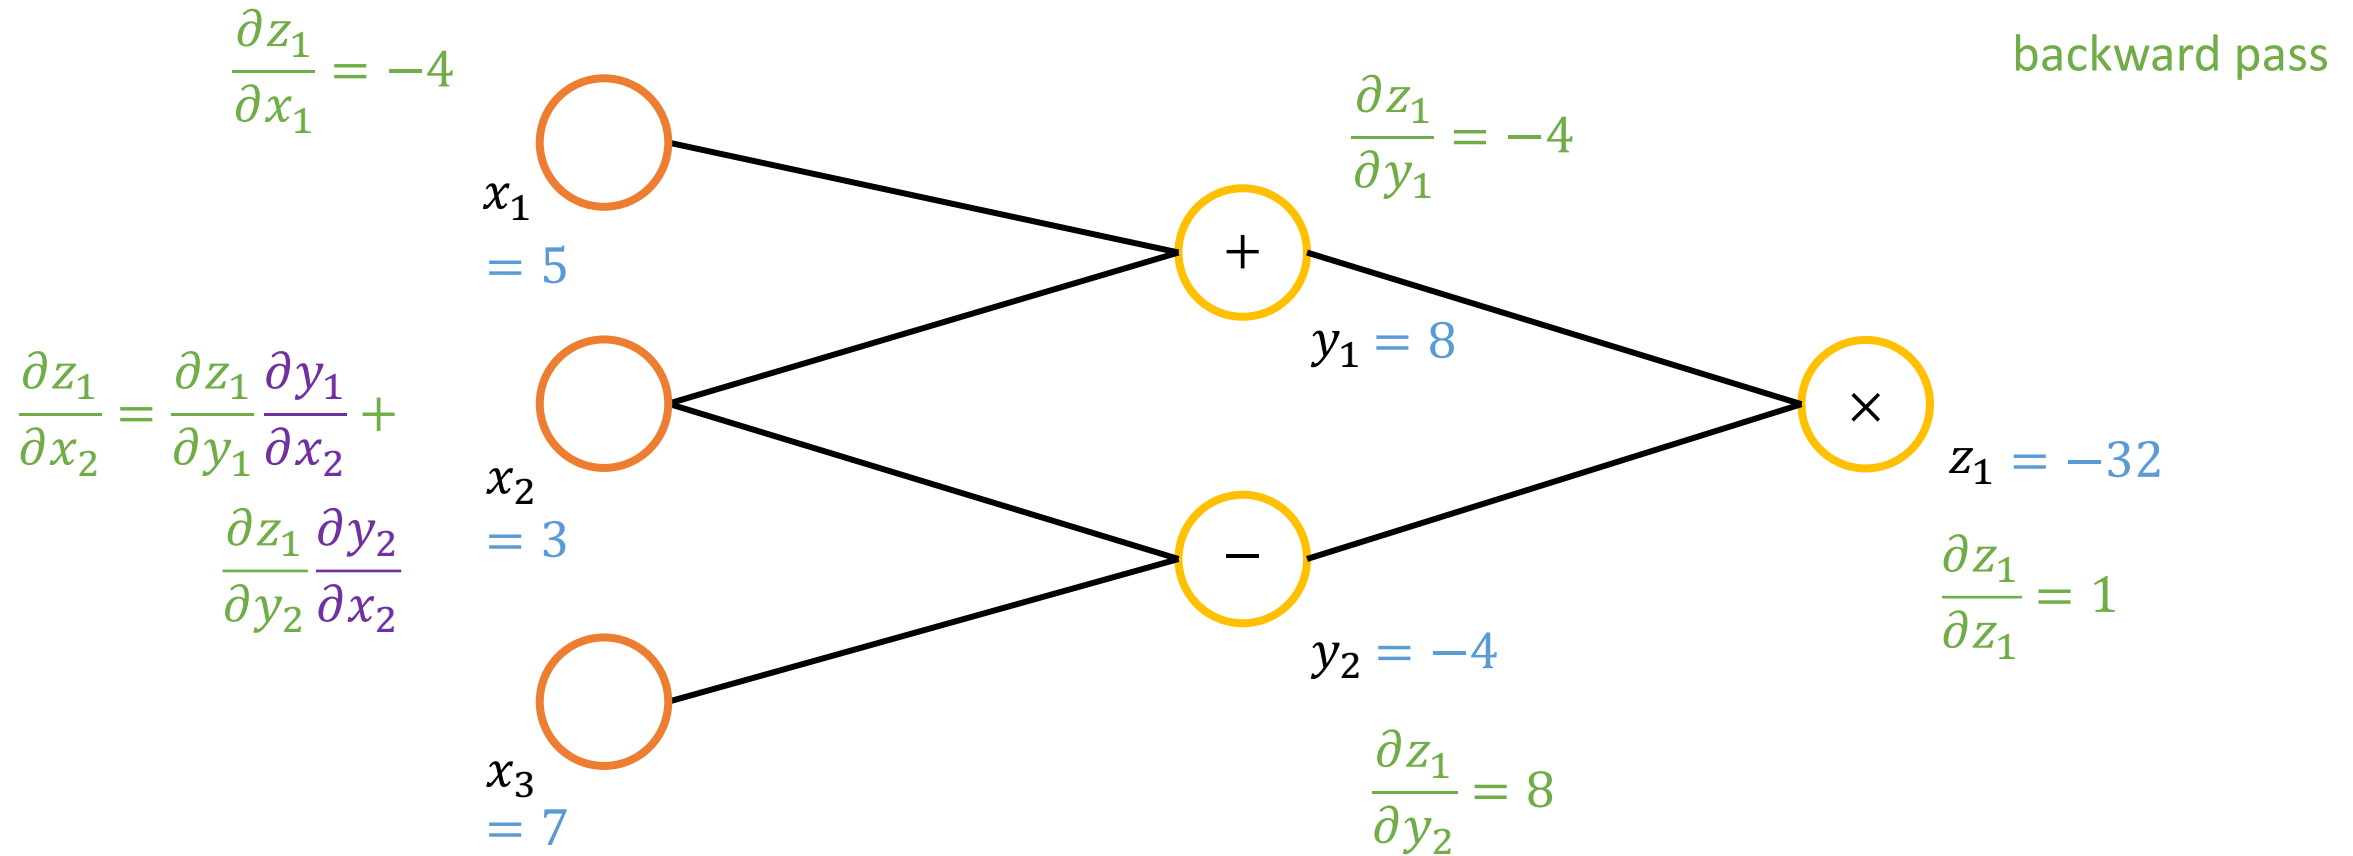
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



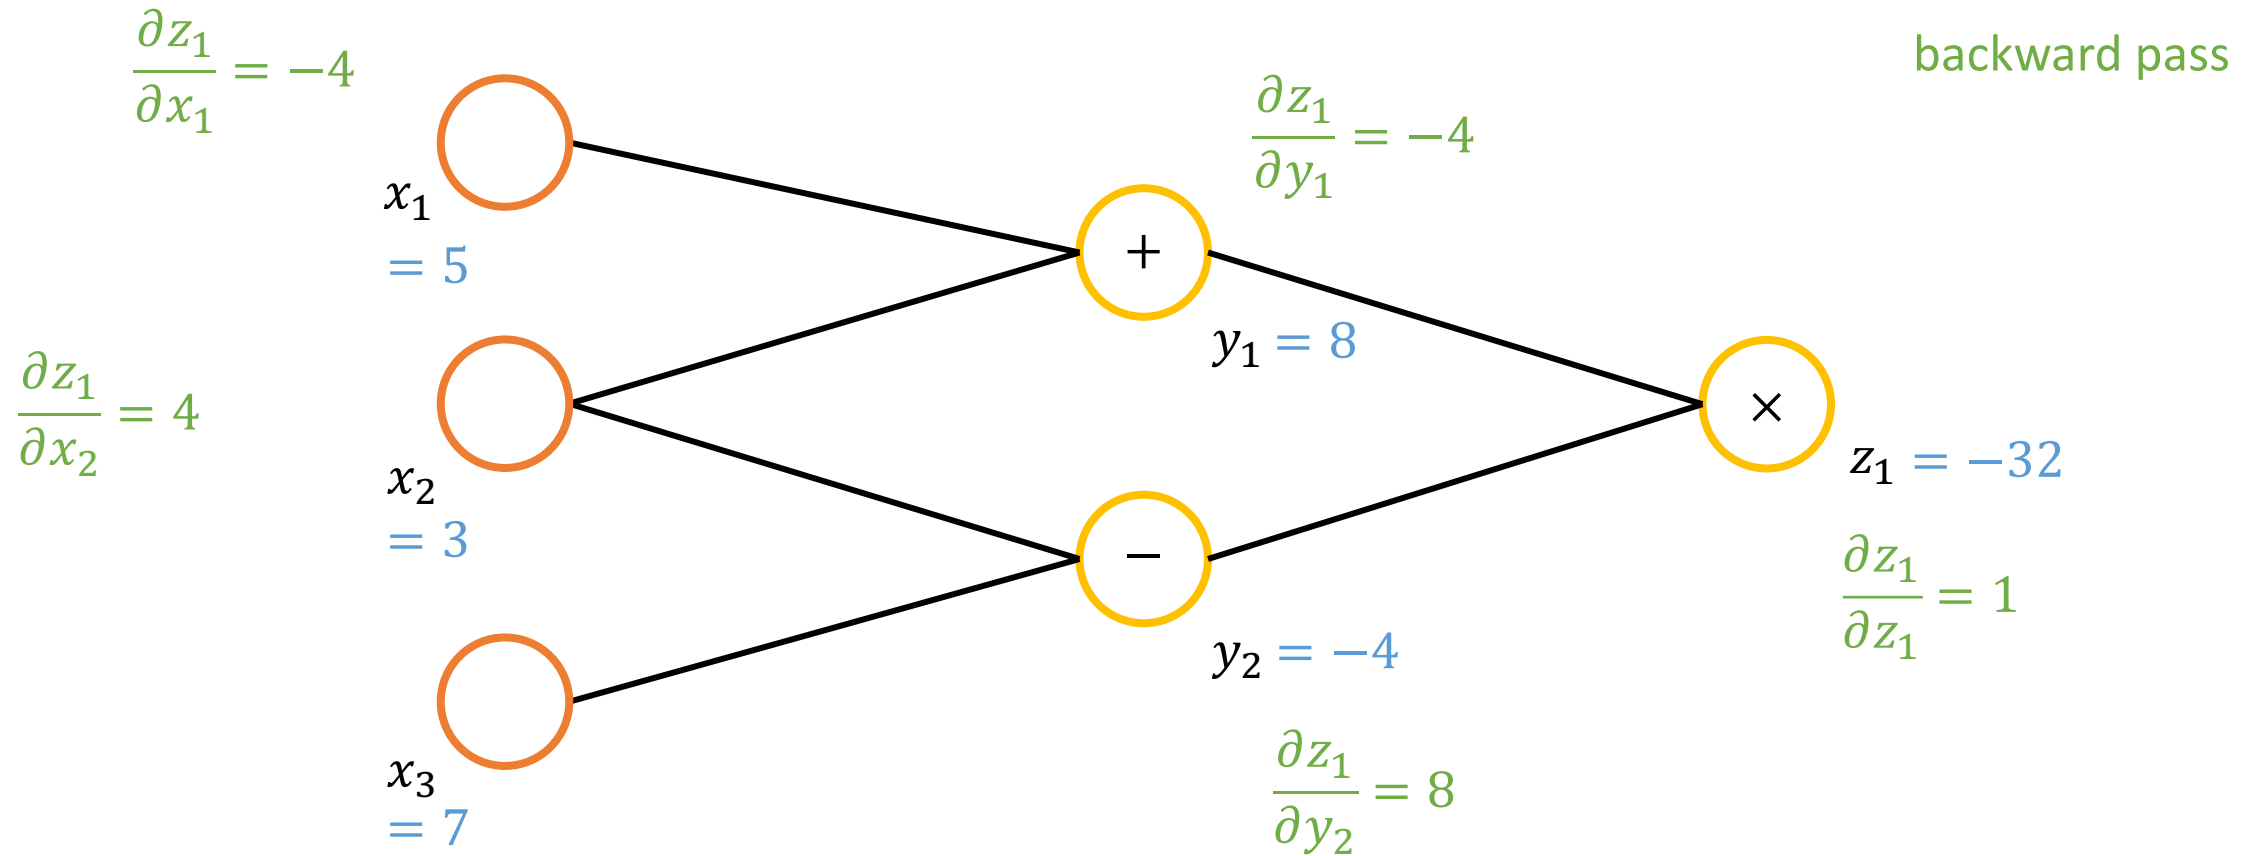
$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1 x_2 - x_1 x_3 + x_2^2 - x_2 x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

Computational graphs, an example



$$z_1 = y_1 \times y_2 = (x_1 + x_2) \times (x_2 - x_3) = x_1x_2 - x_1x_3 + x_2^2 - x_2x_3$$

$$\frac{\partial z_1}{\partial x_1} = x_2 - x_3$$


$$\frac{\partial z_1}{\partial x_2} = x_1 + 2x_2 - x_3$$

$$\frac{\partial z_1}{\partial x_3} = -x_1 - x_2$$

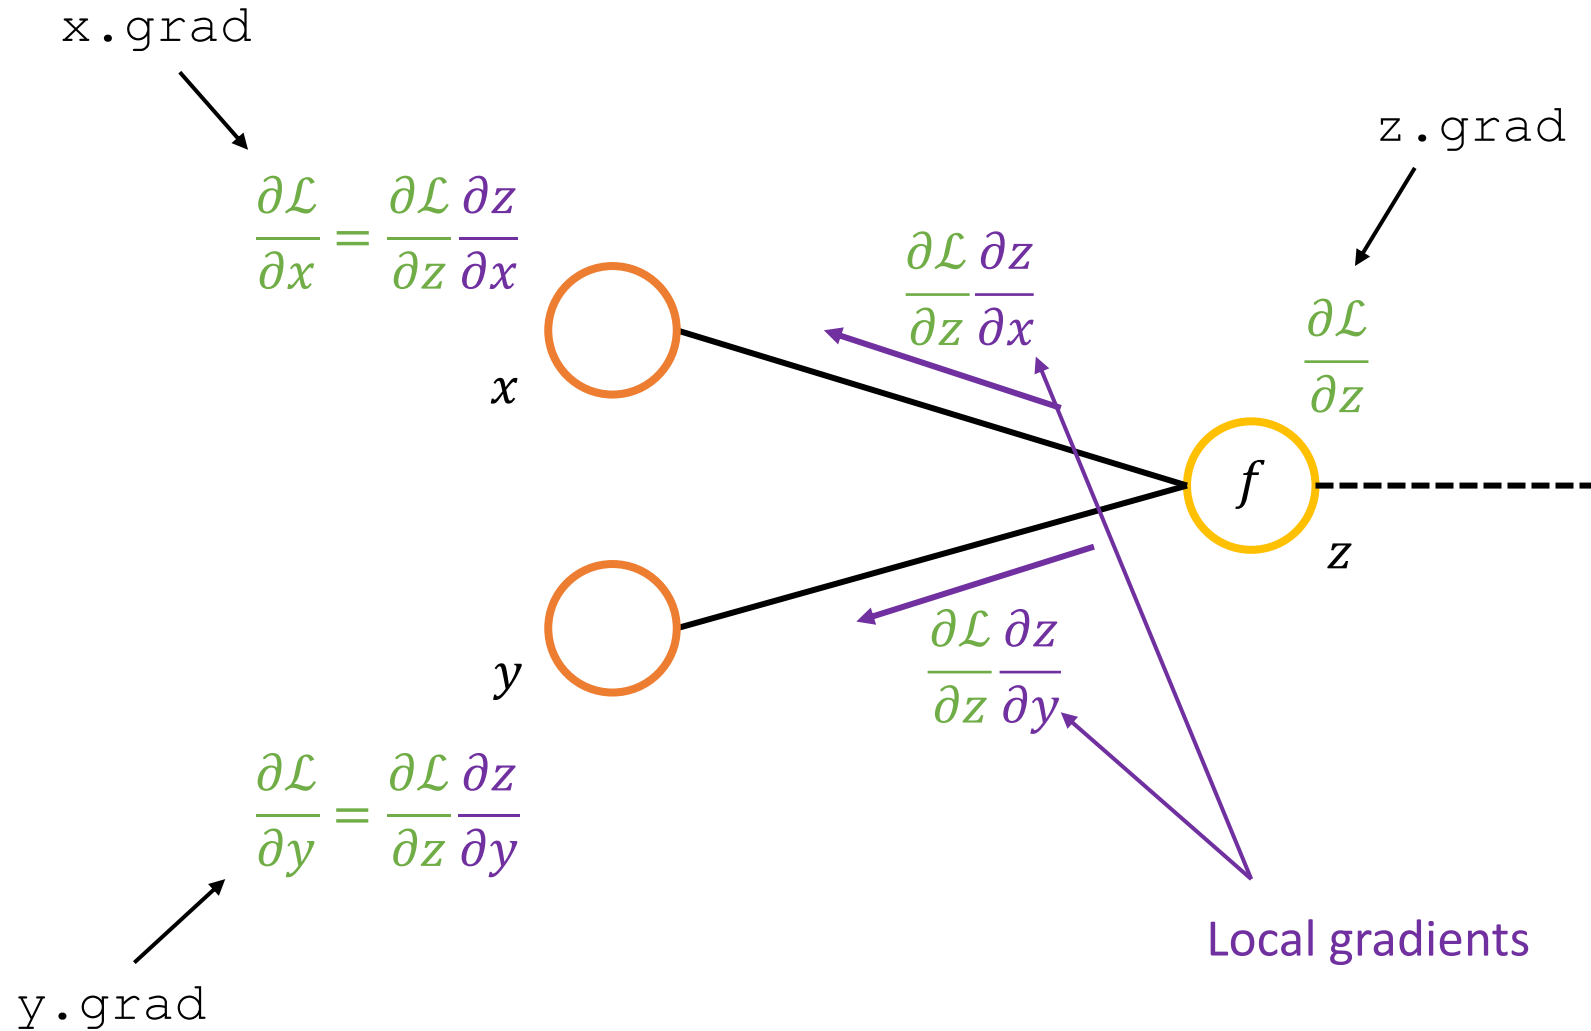
Computational graphs

- The input variables, the learnable parameters are nodes
- The results of elementary operations between nodes are nodes
- Nodes store
 - `node.value`, the result of the operation, computed in the **forward pass**
 - `node.grad`, the gradient of the final output w.r.t. the node variable (at `node.value`), computed in the **backward pass**
- **Graphs** are generally **dynamic**
 - No “fixed”, pre-defined graph that the forward pass is forced to follow
 - The forward pass is defined implicitly by a sequence of operations
 - A backward graph is created on-the-fly during the forward pass (nodes are added to the graph with each operation) so that the gradients can be backpropagated during the backward pass

Backpropagation




Backpropagation



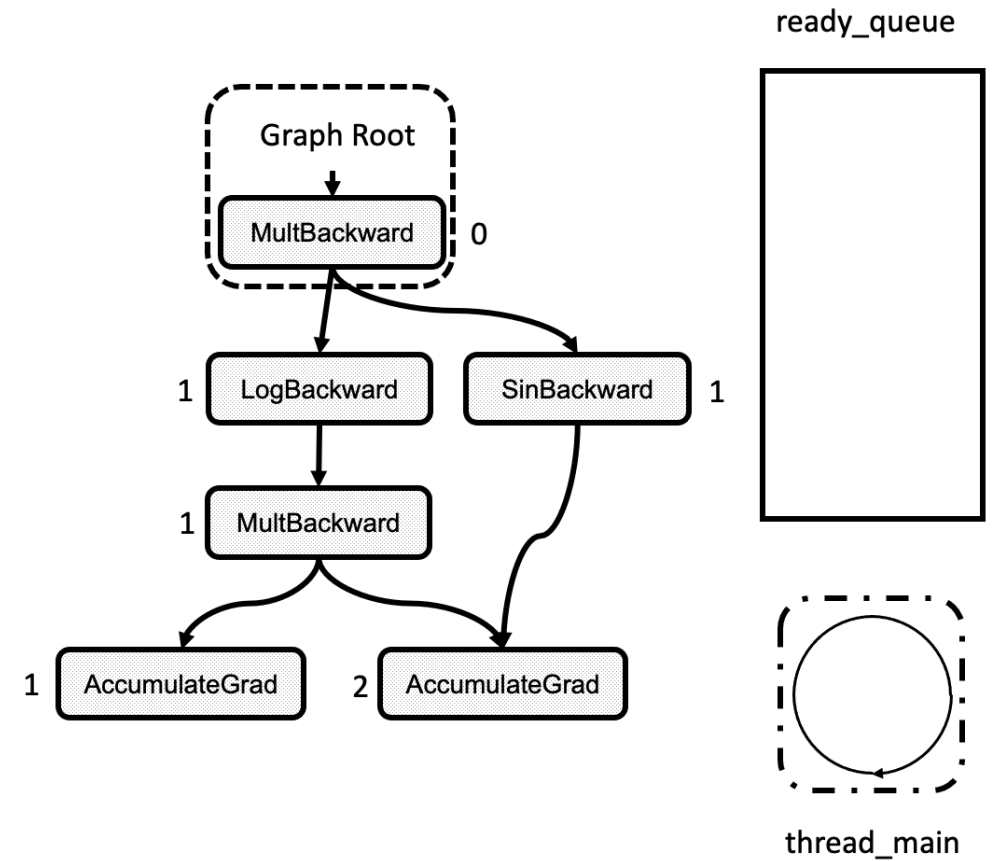
Computational graph execution

Forward pass




Backward pass


(animated source: pytorch.org)




Another example




Another example



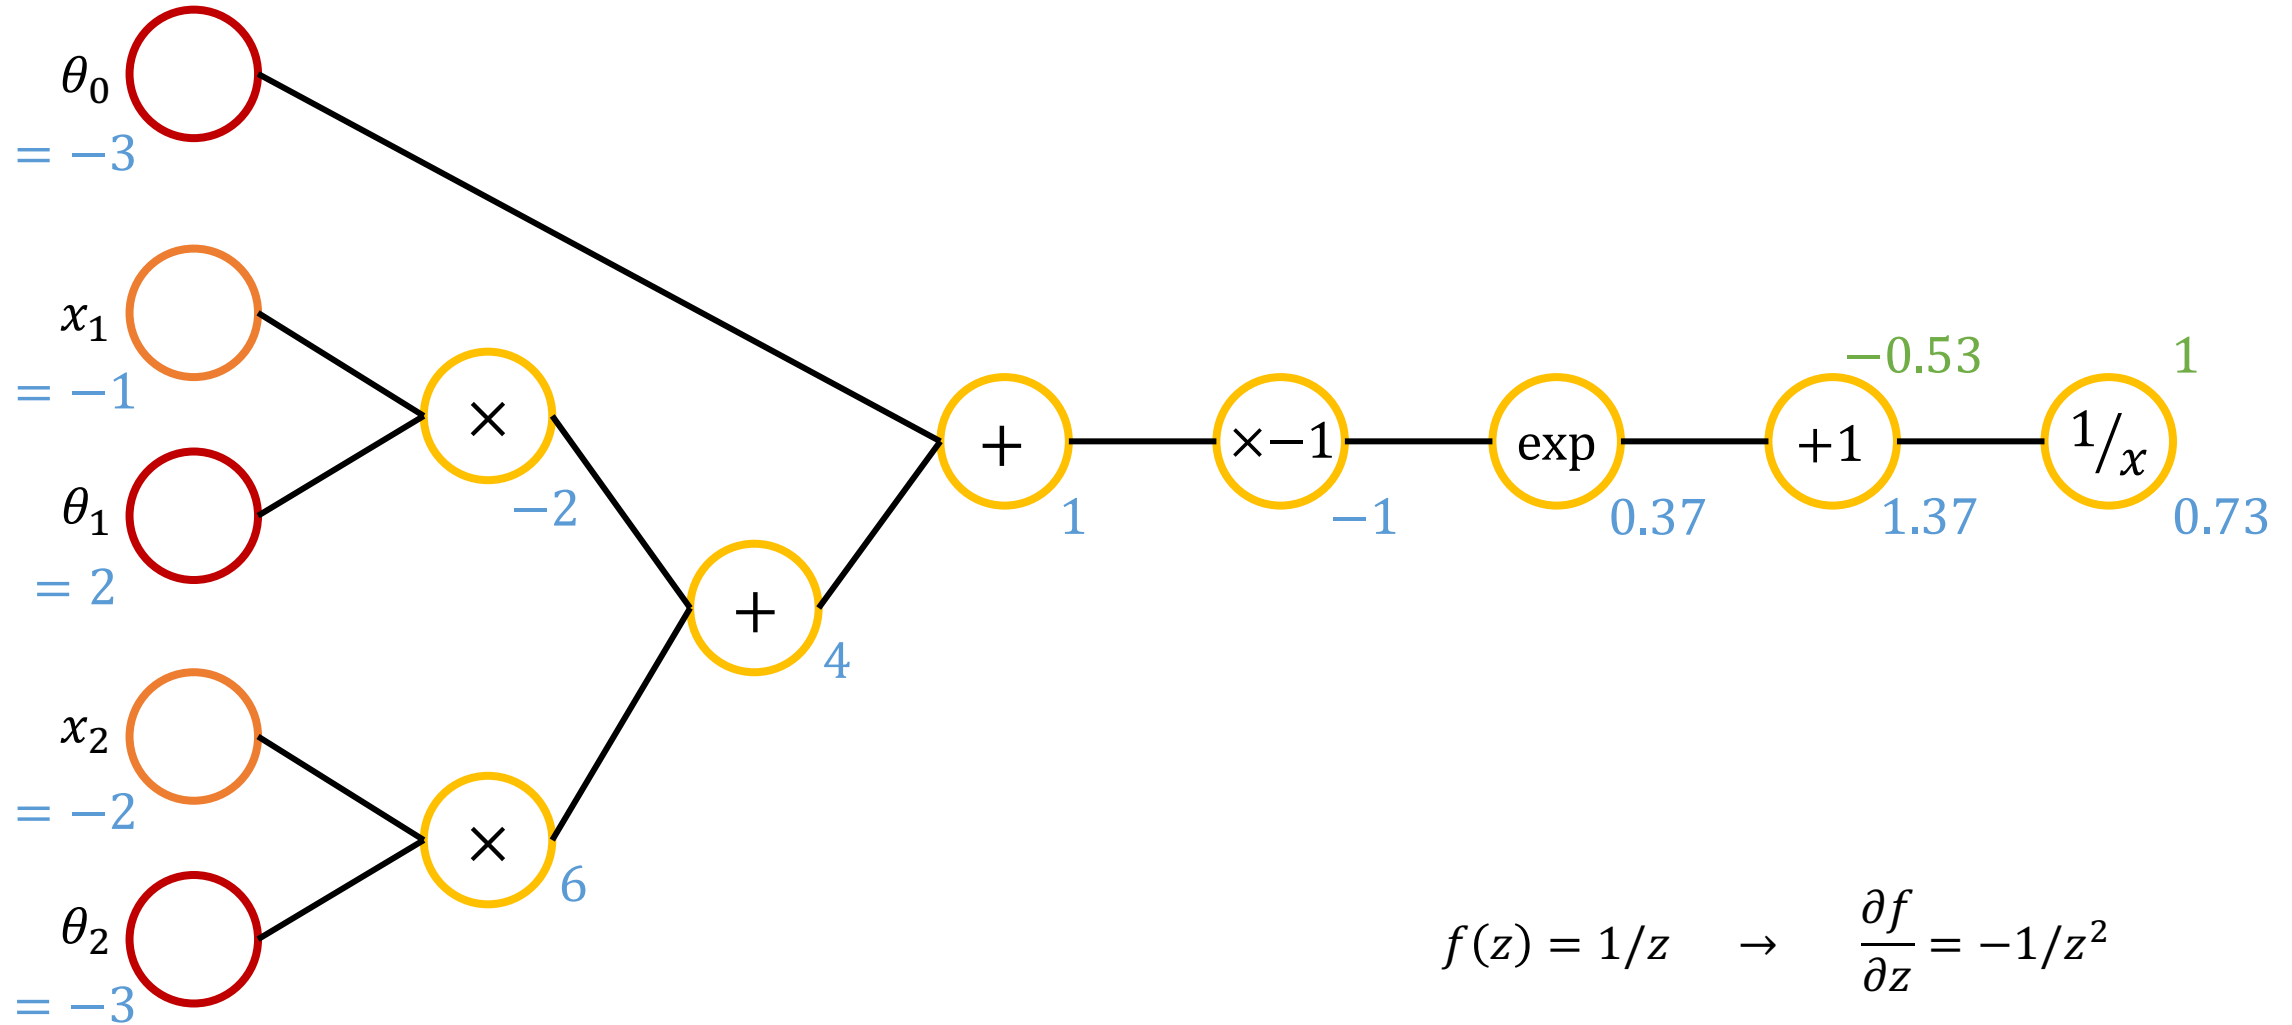
Another example



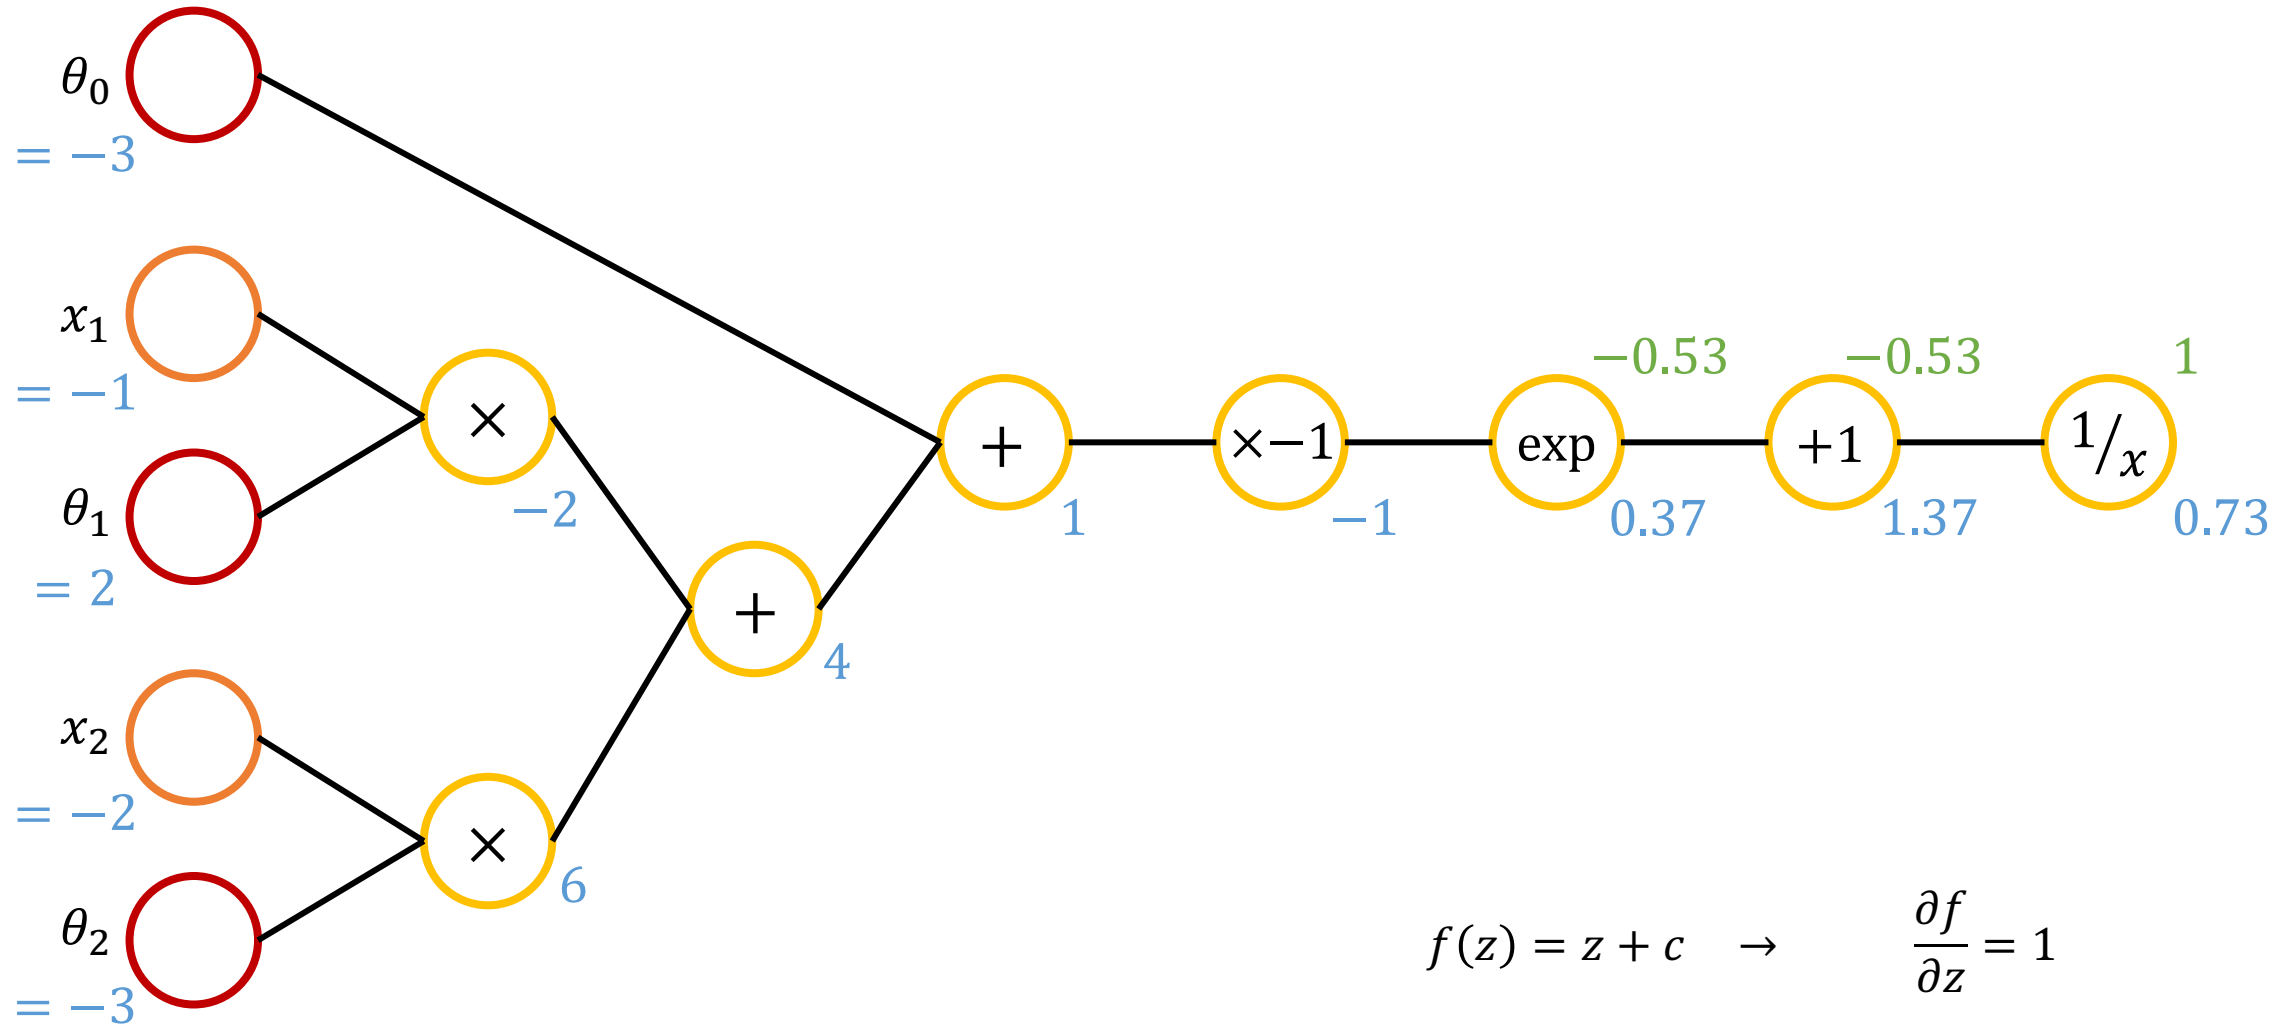
Another example



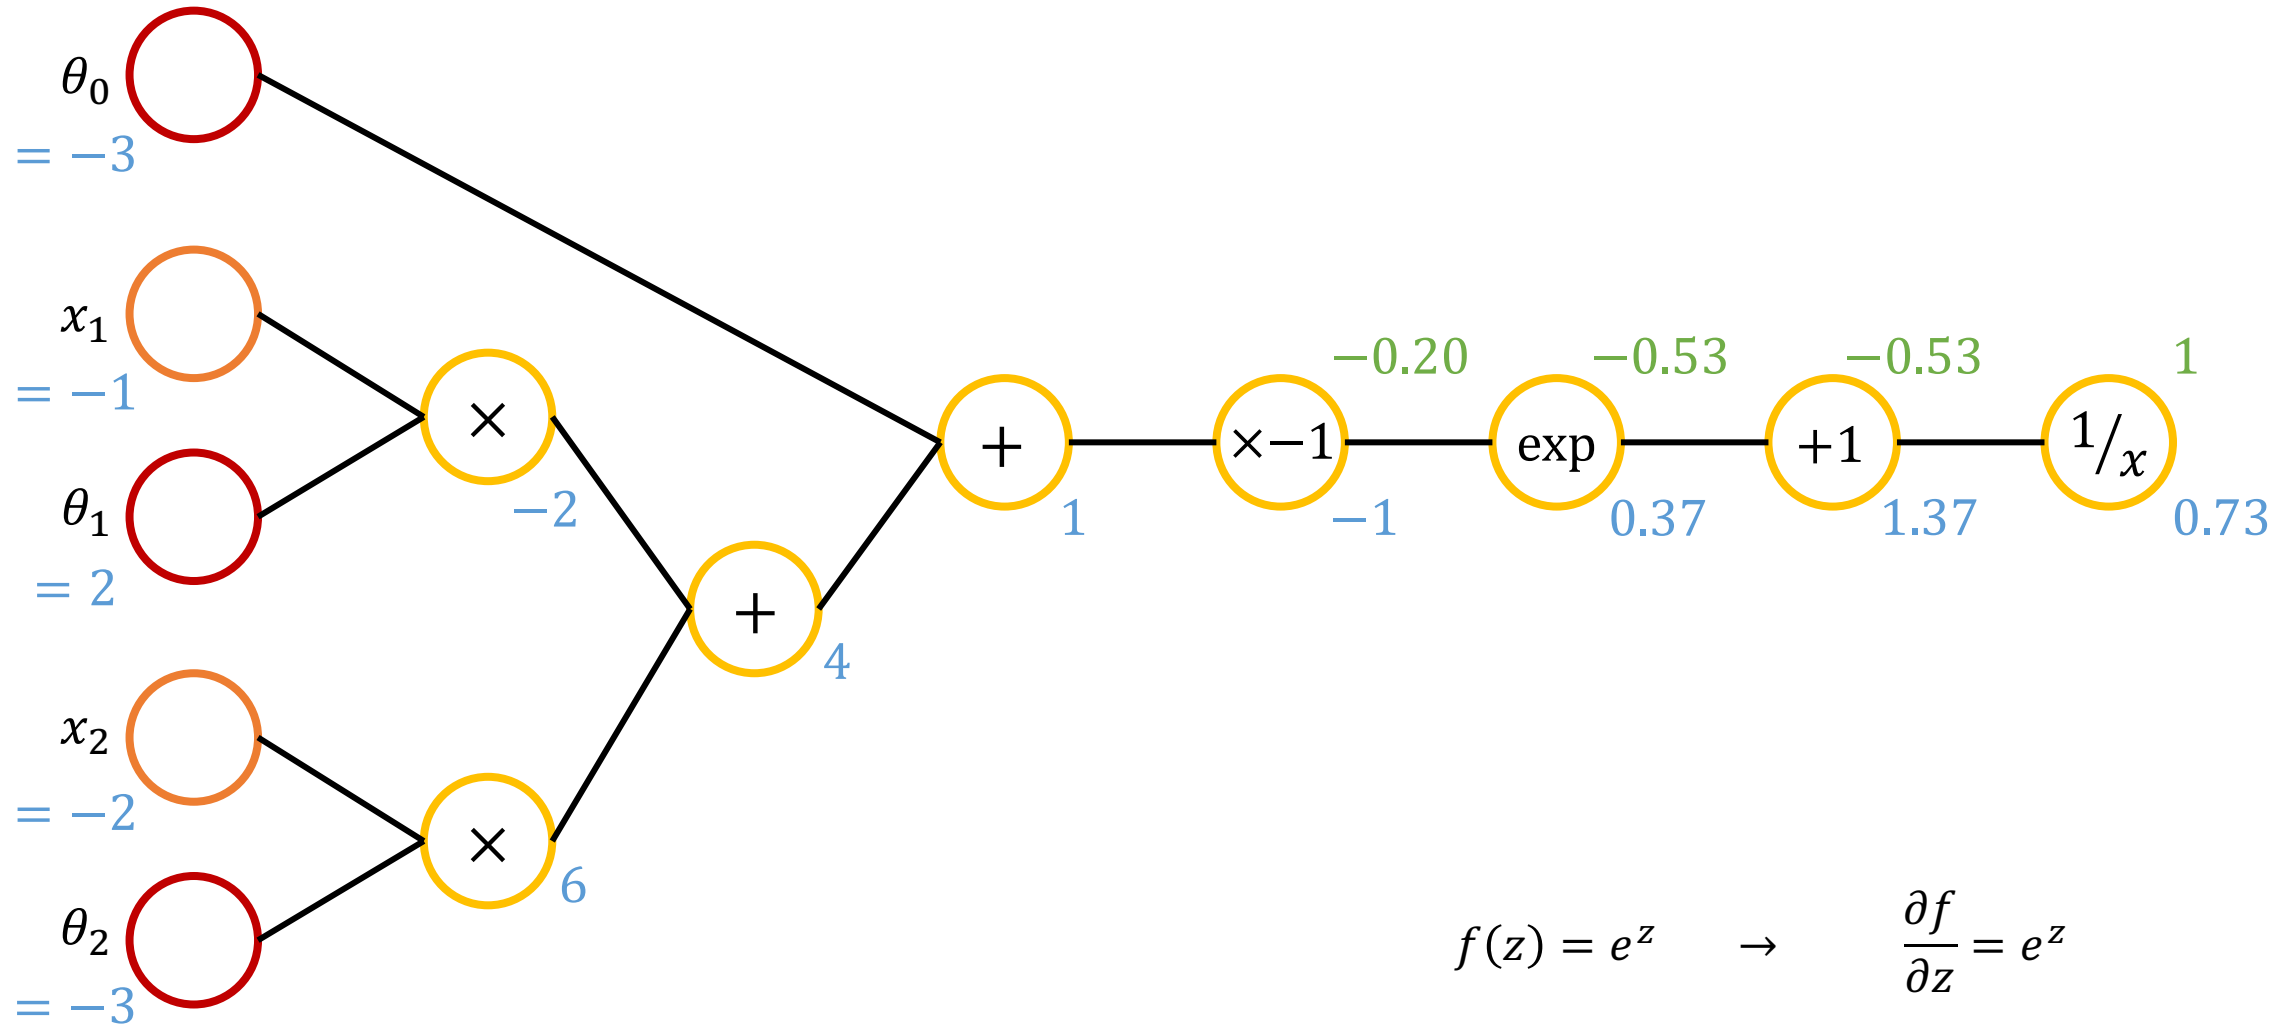
Another example




Another example




Another example



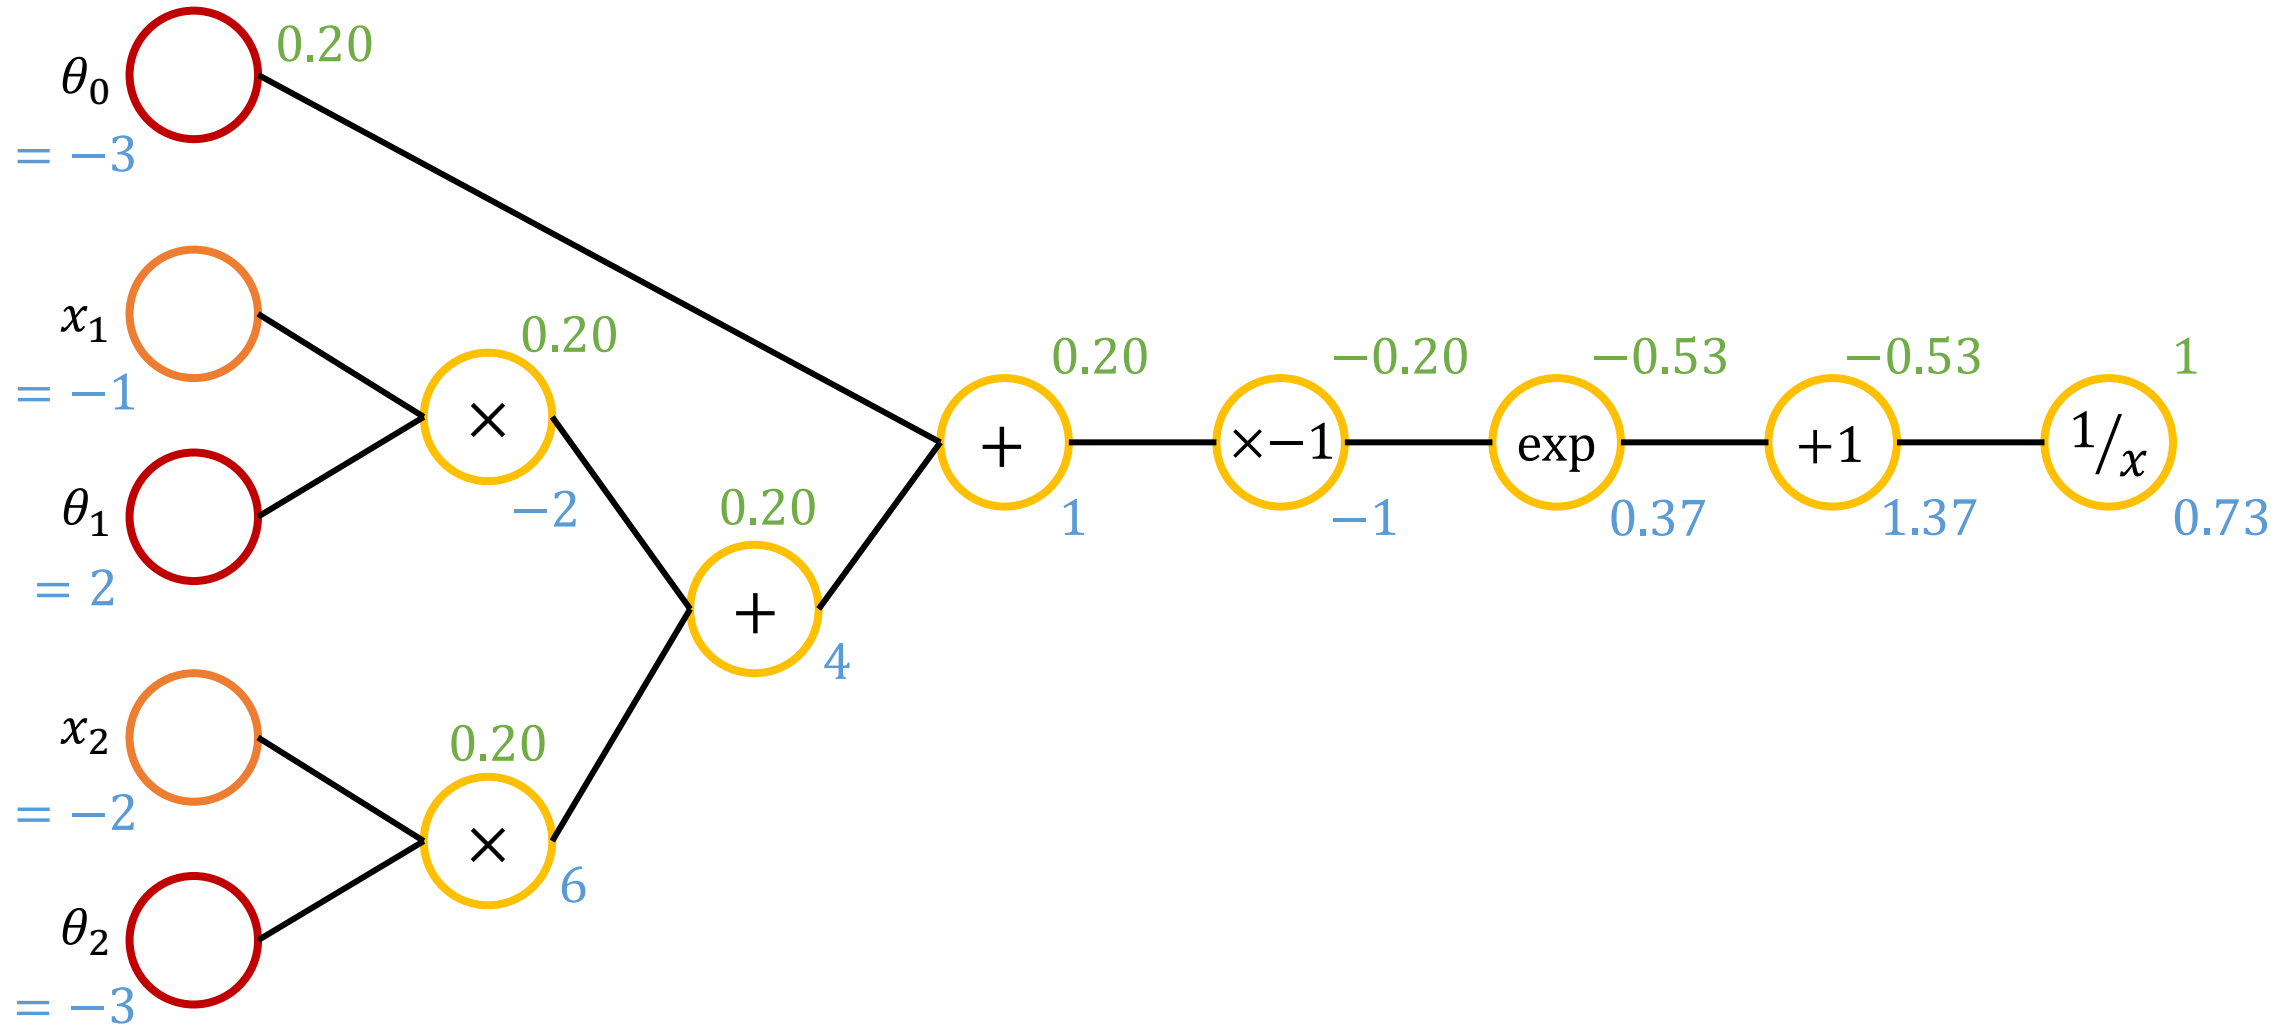
Another example




Another example



Another example



Another example



Patterns in the backward flow

Add gate


→ gradient distributor

Max gate

→ gradient router


Mul gate

→ gradient switcher



Patterns in the backward flow

Gradients add at branches



Back-propagation

Generic optimization algorithm:

1. Forward pass: make a prediction and measure the error
2. Backward pass: Go through each layer in reverse to compute the gradient of the error w.r.t. each variable
3. Update model parameters to reduce the error (gradient descent step)




Letter | Published: 09 October 1986

Learning representations by back-propagating errors

David E. Rumelhart, Geoffrey E. Hinton & Ronald J. Williams

Nature **323**, 533–536 (09 October 1986) | [Download Citation](#)

Back-propagation



Optimizing Neural networks: stochastic gradient descent and variants

(Full-batch) gradient descent is impractical

Remember the loss function:

$$\mathcal{L}(\boldsymbol{\theta}) \triangleq \frac{1}{N} \sum_{i=1}^N \mathcal{L}(y^{(i)}, h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}))$$

- Computing the gradient of $\mathcal{L}(\boldsymbol{\theta})$ is linear in the number of data points N in the batch
- It is impractical to evaluate it frequently for large datasets (e.g., ImageNet has more than $14M$ annotated examples) or datasets for which evaluation of the loss for a single training sample is costly

Solution: **Stochastic Gradient Descent** (SGD), where the gradient $\nabla_{\boldsymbol{\theta}} \mathcal{L}$ is replaced by a minibatch estimate $\hat{\nabla}_{\boldsymbol{\theta}} \mathcal{L}$ computed from $n \ll N$ training samples

Stochastic Gradient Descent

Initialize parameters θ and learning rate η

Repeat until stopping criterion:

Sample a minibatch of n examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}\}$ from the training set with labels $\{y^{(1)}, \dots, y^{(n)}\}$

Compute the minibatch loss function $\tilde{\mathcal{L}}(\theta) \triangleq \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{(i)}, h_{\theta}(\mathbf{x}^{(i)}))$ at θ^t

Compute the gradient of the minibatch loss $\hat{\nabla}_{\theta} \triangleq \nabla_{\theta} \tilde{\mathcal{L}}$ as an estimate of the full-batch gradient


Apply the parameter update: $\theta^{t+1} := \theta^t - \eta \hat{\nabla}_{\theta}$

Epochs

- Each **iteration** ends with an SGD parameter **update step**
- An iteration processes a single minibatch
- The training process often sees 10s of thousands of iterations, so that the full data set is processed several times during optimization
- After a number of iterations, the **full training dataset has been seen** once and exactly once
 - This is called an **epoch**
- At the end of an epoch is generally a good time to evaluate a validation loss on a validation dataset, to evaluate the training stage (are we still making progress in the training?)

Learning rate scheduler

- SGD is no longer guaranteed to converge to an exact local/global minimum
- It may keep overshooting and hover in the vicinity of a minimum
- When we slowly decrease the learning rate η , SGD shows the same convergence behaviour as batch gradient descent, almost certainly converging to a local minimum (non-convex optimization) or global minimum (convex optimization)
 - But large learning rates can have a regularization effect
- The learning rate (LR) can be **adjusted over iterations/epochs**
- Many **learning rate schedulers** have been devised:
 - Simple annealing: linear LR, exponential LR – to reduce the LR as we get closer to convergence
 - Adaptive LR: e.g., reduce the LR by an order of magnitude when a metric plateaus
 - Cyclic LR: triangular, cosine schedules – varies between high LR to escape from current local minimum and low LR to facilitate convergence
 - Cosine annealing with warm restarts: warm restarts facilitate moving to a different basin of attraction
 - 1cycle LR



Huang et al. ICLR 2017

Minibatch size



- Large **minibatch sizes** give a more accurate estimate $\hat{\nabla}_{\theta}$ of $\nabla_{\theta}\mathcal{L}$ but the **returns are not linear**
 - Standard error of $\bar{X} \triangleq \frac{1}{n} \sum_i X_i$ from μ is σ/\sqrt{n} , for n i.i.d. random variables X_i with mean μ and variance σ^2
 - Error decreased by a factor of **10** if number of samples increased by a factor of **100**
- Computing architectures (GPUs...) are underutilized by extremely small batch sizes
→ motivates minimum minibatch size
- Memory usage grows linearly with the minibatch size
→ constrains maximum minibatch size
- Small minibatch sizes have a regularization effect

SGD variants

Variants have been designed to address **SGD defects**:

- “noisy” minibatch gradients that point away from the slope of the true loss function
- **Gradient descent direction suboptimal** in case of highly curved/anisotropic loss profiles

Good resource with the rationale for most variants: <https://www.ruder.io/optimizing-gradient-descent/>



SGD with momentum

Initialize parameters θ , learning rate η , momentum \mathbf{v} and momentum parameter γ

Repeat until stopping criterion:


Sample a minibatch of n examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}\}$ from the training set with labels $\{y^{(1)}, \dots, y^{(n)}\}$

Compute the minibatch loss function $\tilde{\mathcal{L}}(\theta) \triangleq \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{(i)}, h_{\theta}(\mathbf{x}^{(i)}))$ at θ^t

Compute the gradient of the minibatch loss $\hat{\nabla}_{\theta} \triangleq \nabla_{\theta} \tilde{\mathcal{L}}$

Compute momentum update $\mathbf{v}^{t+1} := \gamma \mathbf{v}^t + \eta \hat{\nabla}_{\theta}$

Apply the parameter update: $\theta^{t+1} := \theta^t - \mathbf{v}^{t+1}$



SGD with Nesterov momentum: reduce overshooting

Initialize parameters θ , learning rate η , momentum v and momentum parameter γ

Repeat until stopping criterion:

Sample a minibatch of n examples $\{x^{(1)}, \dots, x^{(n)}\}$ from the training set with labels $\{y^{(1)}, \dots, y^{(n)}\}$


Apply interim update $\tilde{\theta} := \theta^t - \gamma v^t$

Compute the minibatch loss function $\tilde{\mathcal{L}}(\theta) \triangleq \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{(i)}, h_{\theta}(x^{(i)}))$ at $\tilde{\theta}$

Compute the gradient of the minibatch loss $\hat{\nabla}_{\theta} \triangleq \nabla_{\theta} \tilde{\mathcal{L}}$ at $\tilde{\theta}$

Compute momentum update $v^{t+1} := \gamma v^t + \eta \hat{\nabla}_{\theta}$

Apply the parameter update: $\theta^{t+1} := \theta^t - v^{t+1}$



RMSprop: normalize gradients by a moving average

Initialize parameters θ , learning rate η , squared gradient moving average $E[g^2]$, parameter $0 < \gamma < 1$ and small constant ϵ

Repeat until stopping criterion:

Sample a minibatch of n examples $\{x^{(1)}, \dots, x^{(n)}\}$ from the training set with labels $\{y^{(1)}, \dots, y^{(n)}\}$

Compute the minibatch loss function $\tilde{\mathcal{L}}(\theta) \triangleq \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{(i)}, h_{\theta}(x^{(i)}))$ at θ^t

Compute the gradient of the minibatch loss $\hat{V}_{\theta} \triangleq \nabla_{\theta} \tilde{\mathcal{L}}$ at θ^t

Update running average $E[g^2]_t := \gamma E[g^2]_{t-1} + (1 - \gamma) \hat{V}_{\theta} \odot \hat{V}_{\theta}$

Compute the parameter update $\Delta \theta^t := - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} \odot \hat{V}_{\theta}$

Apply the parameter update: $\theta^{t+1} := \theta^t + \Delta \theta^t$

Adam: use moving average of gradient

Initialize parameters θ , learning rate η , gradient moving average $\mathbf{m} := \mathbf{0}$, squared gradient moving average $\mathbf{v} := \mathbf{0}$, parameters $0 < \beta_1, \beta_2 < 1$ and small constant ϵ

Repeat until stopping criterion:

Sample a minibatch of n examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}\}$ from the training set with labels $\{y^{(1)}, \dots, y^{(n)}\}$

Compute the minibatch loss function $\tilde{\mathcal{L}}(\theta) \triangleq \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y^{(i)}, h_{\theta}(\mathbf{x}^{(i)}))$ at θ^t

Compute the gradient of the minibatch loss $\hat{\nabla}_{\theta} \triangleq \nabla_{\theta} \tilde{\mathcal{L}}$ at θ^t


Update running averages: $\mathbf{m}_t := \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \hat{\nabla}_{\theta}$ and $\mathbf{v}_t := \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \hat{\nabla}_{\theta} \odot \hat{\nabla}_{\theta}$

De-bias the moment estimates: $\hat{\mathbf{m}}_t := \mathbf{m}_t / (1 - \beta_1^t)$ and $\hat{\mathbf{v}}_t := \mathbf{v}_t / (1 - \beta_2^t)$

Compute the parameter update $\Delta \theta^t := -\frac{\eta}{\sqrt{\hat{\mathbf{v}}_t + \epsilon}} \odot \hat{\mathbf{m}}_t$

Apply the parameter update: $\theta^{t+1} := \theta^t + \Delta \theta^t$

Choice of learning rate



Regularization

Regularization

Regularization is useful when the optimization **problem is underdetermined**


- This is common in ML and leads to overfitting

Techniques **to avoid overfitting** applicable with NNs:

- Early stopping
- L^p regularization
- Dropout
- Data augmentation

Early stopping

- **Interrupt the training** when the **validation loss** stops decreasing for a few epochs
- Keep the best result (lowest validation loss)
- Necessitates to monitor the loss on a validation dataset at the end of each epoch (good practice anyway)



L^p regularization

Adds another term to the loss function in addition to the data term $\mathcal{L}_{\text{data}}(\boldsymbol{\theta}) \triangleq \frac{1}{N} \sum_{i=1}^N \mathcal{L}(y^{(i)}, h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}))$ seen in the previous slides

L^2 regularization

a.k.a. Tikhonov regularization, quadratic regularization

$$\mathcal{L} \triangleq \mathcal{L}_{\text{data}} + \lambda \|\boldsymbol{\theta}\|_2^2$$

Can be implemented directly via the optimizer under the name “weight decay”

L^1 regularization

$$\mathcal{L} \triangleq \mathcal{L}_{\text{data}} + \lambda \|\boldsymbol{\theta}\|_1$$

Weight decay


- Weight decay is often taken as synonymous to L2 regularization
- In the original work of Hanson et Pratt (1988), weight decay corresponds to the red term in the following parameter update:

$$\Delta\theta_t := -\lambda\theta_t - \alpha\nabla\mathcal{L}_{\text{data}}(\theta_t)$$

- When using SGD, this is indeed identical to L2 regularization with a weight λ/α
- But popular modern optimizers (RMSprop, Adam) use adaptive updates with modified gradients
- In this case, it is different to:
 - Use L2 regularization, or equivalently add $+\lambda\theta_t$ to $\nabla\mathcal{L}_{\text{data}}(\theta_t)$ before computing the normalized updates of RMSprop / Adam
 - Use weight decay i.e., add the term $-\lambda\theta_t$ in the final computation of $\Delta\theta_t$
- The latter strategy is shown in Loshchilov et Hutter, ICLR 2019, to be superior to the former strategy:
 - a. better generalization error
 - b. broader optimal parameter sets that better decouple optimal learning rate and weight decay values
- It is implemented in the **AdamW optimizer** (“Adam with decoupled weight decay”)
 - Careful: in RMSprop and Adam, “weight decay” actually refers to the former strategy (L2 regularization)

Dropout

- At each iteration, **every neuron** (input or hidden), for each data point in the minibatch, **has a probability p of being “dropped out”**: its value is replaced by 0
- Prevents co-adaptation of neurons



Source: medium.com


- p is called dropout rate
- can be set to 0.5 or smaller values
- At test time, neurons are not dropped out
 - Apply a small correction to correct for the fact that each neuron on average sees input from more neurons
- Dropout **implemented as a layer** in the network, to be placed right after the layer to which it applies
- Very effective for fully connected layers

Data augmentation

Generate new training samples with artificial variability **from existing training samples**, artificially boosting the size of the dataset

Applied random transformations need to be realistic

- Intensity transformations (contrast, color, noise, etc.)
- Spatial transformations: crop, scale, rotate, shift (translate), flip




Implementation, etc.

Weight initialization

How to initialize weights $\theta_{i,j}^l$?

- ~~Constant value $\theta_{i,j}^l := 0$~~ does not break symmetry \rightarrow bad idea
- Small random value: $\theta_{i,j}^l := 0.01 \cdot \mathcal{N}(0,1)$ or $\theta_{i,j}^l := \mathcal{N}(0,1)/\sqrt{M_l}$ or $\theta_{i,j}^l := \sqrt{2/M_l} \cdot \mathcal{N}(0,1)$ (He et al., 2015)
 - Scaling by number of neurons M_l input to the layer allows to keep variance of output constant
- “Xavier” initialization (Gloriot and Bengio, 2010): $\theta_{i,j}^k = \mathcal{U}(-1/\sqrt{n}, +1/\sqrt{n})$ where $\mathcal{U}(-a, +a)$ is the uniform distribution in the interval delimited by $-a$ and $+a$
- ...

Normalizing input data



Normalizing intermediate activations

- Normalizing the input layer is great, but
- What about the intermediate activations?
- The inputs to intermediate layers are not normalized since different neurons have different weights and different activation statistics
- All previous layers $1 \dots (l - 1)$ have an effect on the l th layer's activations
- When backpropagating, these interactions are not accounted for: the weights of the layer l will be updated without accounting for the fact that the previous layers are updated too and their activation statistics change
- So-called **“internal covariate shift”** = change of distributions as network updates

Idea: Normalize the mean and variance in intermediate layers as well

Batch normalization

Let $\mathbf{x}_1, \dots, \mathbf{x}_M$ be intermediate inputs

- C -dimensional quantities if C is the number of neurons in the previous layer
- For us so far, m will index the sample index in the minibatch and $M = n$ is the minibatch size (because we deal with tabular data; for image data, this will be slightly different as per next slides)
- Intermediate inputs can be before (more rarely) or after (more commonly) the nonlinearity

Compute the **normalized intermediate input**:

$$\hat{\mathbf{x}}_m := \frac{\mathbf{x}_m - \boldsymbol{\mu}}{\sqrt{\boldsymbol{\sigma}^2 + \epsilon}}$$

Where:

- $\boldsymbol{\mu} := \frac{1}{M} \sum_m \mathbf{x}_m$
- $\boldsymbol{\sigma}^2 := \frac{1}{M} \sum_m (\mathbf{x}_m - \boldsymbol{\mu})^2$

Output $\mathbf{h}_m := \boldsymbol{\gamma} \odot \hat{\mathbf{x}}_m + \boldsymbol{\beta} \triangleq BN_{\boldsymbol{\gamma}, \boldsymbol{\beta}}(\mathbf{x}_m)$

where $\boldsymbol{\gamma}, \boldsymbol{\beta}$ are learnable parameters

Batch normalization

Let $\mathbf{x}_1, \dots, \mathbf{x}_M$ be intermediate inputs

- C -dimensional quantities if C is the number of neurons in the previous layer
- For us so far, m will index the sample index in the minibatch and $M = n$ is the minibatch size (because we deal with tabular data; for image data, this will be slightly different as per next slides)
- Intermediate inputs can be before (more rarely) or after (more commonly) the nonlinearity

Compute the **normalized intermediate input**:

$$\hat{\mathbf{x}}_m := \frac{\mathbf{x}_m - \boldsymbol{\mu}}{\sqrt{\boldsymbol{\sigma}^2 + \epsilon}}$$

Where:

- $\boldsymbol{\mu} := \frac{1}{M} \sum_m \mathbf{x}_m$
- $\boldsymbol{\sigma}^2 := \frac{1}{M} \sum_m (\mathbf{x}_m - \boldsymbol{\mu})^2$

At test time, minibatch averages are replaced by a fixed exponential moving average (EMA, computed during training)

Output $\mathbf{h}_m := \boldsymbol{\gamma} \odot \hat{\mathbf{x}}_m + \boldsymbol{\beta} \triangleq BN_{\boldsymbol{\gamma}, \boldsymbol{\beta}}(\mathbf{x}_m)$

where $\boldsymbol{\gamma}, \boldsymbol{\beta}$ are learnable parameters

Batch normalization


Downsides of BN:

- Does not work well with small minibatches (estimated statistics are unreliable)
- Uses different statistics for inference / test-time than for training
- Does not always work well with L2 regularization (Hoffer et al. NeuRIPS 2018) or dropout (Li et al. CVPR 2019)
- Introduces a significant computational overhead during training
The EMA also introduces an additional memory requirement


Alternatives to batch normalization

S : signal size ($H \times W$ for images), C : number of features/channels, $|B|$: minibatch size


Blue colors indicate the dimensions over which the averages are computed for the normalization




Batch Normalization




Layer Normalization




Group Normalization




Instance Normalization



Tabular



Images



Sequential